

Presentation of

Sim-Diasca

Simulation of Discrete Systems of All Scales

Friday, April 7, 2023

Contact: Olivier Boudeville
(olivier.boudeville@edf.fr)

More information:

- <http://www.sim-diasca.com>
- <https://github.com/Olivier-Boudeville-EDF/Sim-Diasca>

Major French Utility





Presentation outline: introducing Sim-Diasca

1. All Sim-Diasca in one slide
2. Requirements & technical answer
3. Algorithmic choices
4. Technical design & features
5. What is Sim-Diasca?
 - Functional Service & Key Points
 - Software Architecture
6. Future work
7. Conclusion
8. Appendices

Sim-Diasca

Simulation of Discrete Systems of All Scales

Sim-Diasca is a **concurrent** (parallel and distributed) **generic discrete-time simulation engine** aiming at **maximum scalability** (millions of complex model instances in interaction).

- Generic, domain-agnostic: can be applied to a wide range of large-scale discrete simulation targets, from ecosystems to vast IT infrastructures
- Typically suitable for simulations in the field of Complex Systems
(whereas most of the tools for that are sequential and can hardly scale)
- Sim-Diasca (simulation engine) + models + simulation case(s) ⇒ a simulator
- Fully implemented in a functional language, Erlang (<http://erlang.org>)
- Supported platforms: GNU/Linux (from single laptops to full-blown HPC clusters)
- Used by EDF and third parties, maintained by EDF R&D
- Released since 2010 by EDF R&D as free software (LGPL licence)

Sim-Diasca requirements & technical answer

Functional Requirements:

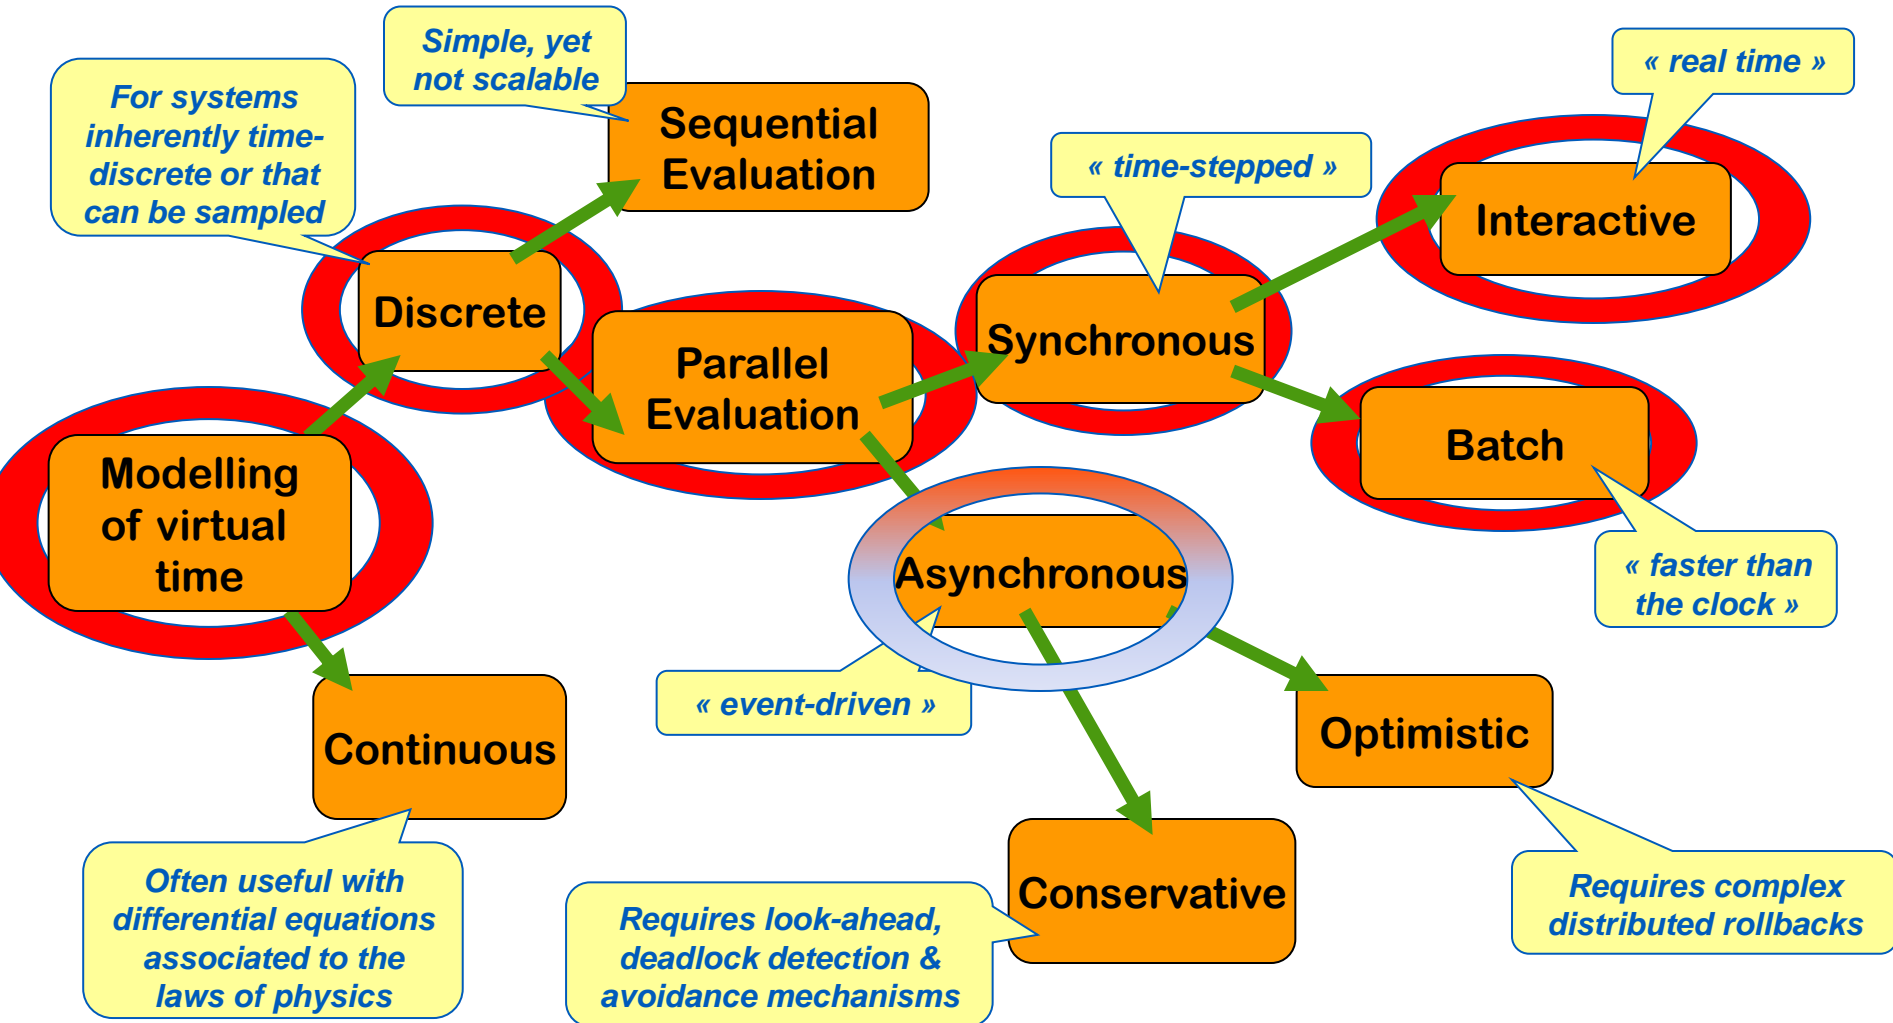
- Multiple **usages** anticipated (simulator/emulator/digital-twin)
- Need for **correctness** in the evaluation of all kinds of models
- Simulation use cases require to be able to:
 - **Replay** at will any given trajectory of the target system
 - **Correlate** directly a change into the simulation results to a change into the inputs
 - **Explore** all possible trajectories of the system, moreover in a fair, representative way
- Need to be able to simulate **very large systems**, potentially involving dozens of millions of interacting actors

Technical Answer:

- **Batch** or **Interactive** mode of operation
- Respect of **causality**
- Support for **stochastic** models
- **Total reproducibility**
- Support of a certain form of “**ergodicity**”, i.e. a guarantee that:
 - All possible outcomes according to the models *can* actually occur in the simulations
 - And that their probability of showing up in the simulations is close to the one that can be deduced from the models
- Ability to **scale up** significantly:
 - At the algorithmic level: maximal parallelization of the evaluation of models
 - At the level of computing resources: harness multicores, SMP, clusters and other High Performance Computing solutions

Sim-Diasca main algorithmic choices

The nodal point is how the simulation time is managed:



► Choices for the Sim-Diasca mode of operation are shown within **red ellipses**.

Sim-Diasca technical design & features

▶ As an answer to the requirements, a simulation engine:

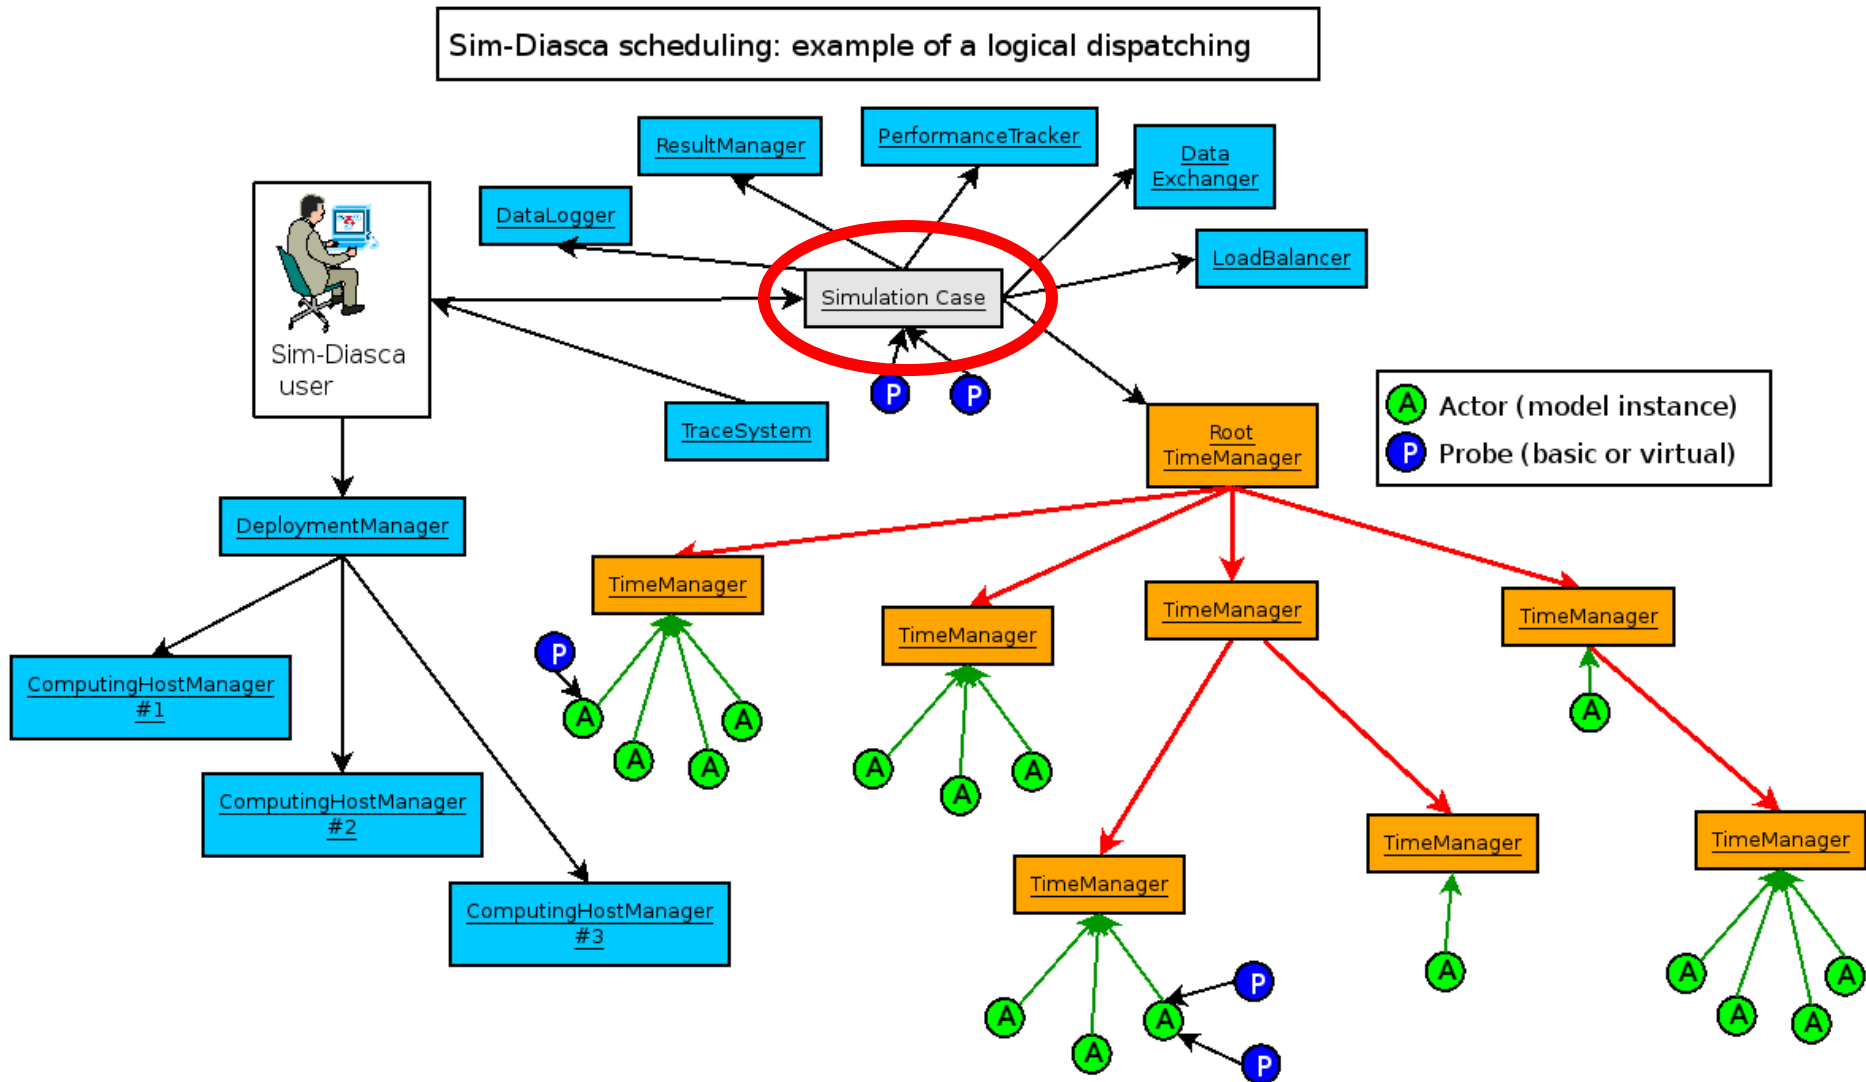
- **Based on discrete events:** the elements of the target system (model instances, a.k.a. actors) exchange messages and update their state accordingly
- **Synchronous** (« *time-slicing* », « *time-stepped* »):
 - A fundamental simulation frequency is defined (by default, 50Hz)
 - In interactive mode, the engine adjusts its time steps (ticks) to the real (wall-clock) time
 - In batch mode, the engine processes its ticks at maximum speed, and jumps automatically over periods without any possible activity of actors (quite similarly to asynchronous approaches)
- **Intensely concurrent:**
 - Distributed simulation: a single simulation can spread over a set of computing hosts (e.g. HPC cluster)
 - and Parallel, i.e. taking advantage, for each computing host, of all cores of all processors
 - The algorithm allows, at each scheduled logical moment, to evaluate all model instances in parallel!
 - Scalability-wise, at the end of 2010, the threshold of 1 million instances of rather complex models could be reached
- **Granting a large freedom and expressivity to models:**
 - Modelling: object-oriented approach, whence implementation directly derives, based on a concurrent high-level functional language (Erlang)
 - Flexible and powerful scheduling policies for actors (fully passive, periodical, or driving their own behaviour arbitrarily)
 - Stochastic support: any number of stochastic variables per actor, respecting built-in or model-defined probability density functions
 - Very few constraints apply to models (e.g. no pre-established fan-in/fan-out, no look-ahead needed); no causality-induced time biases (as many logical moments - « diasca » - as needed will be created in a tick to sort out causality)
- **Providing all needed features to build easily a simulator:**
 - Engine features: automatic and parallel deployment (code and data), management of results, load-balancing, integration to most platforms (e.g. clusters), distributed trace system, tuning for scalability, etc.
 - Model features: automatic message reordering, stochastic support, probe support (autonomous and database-based)

And only the relevant ones

▶ Central point: the massive parallelism is achieved without prejudice to the targeted simulation properties

Functional Services:

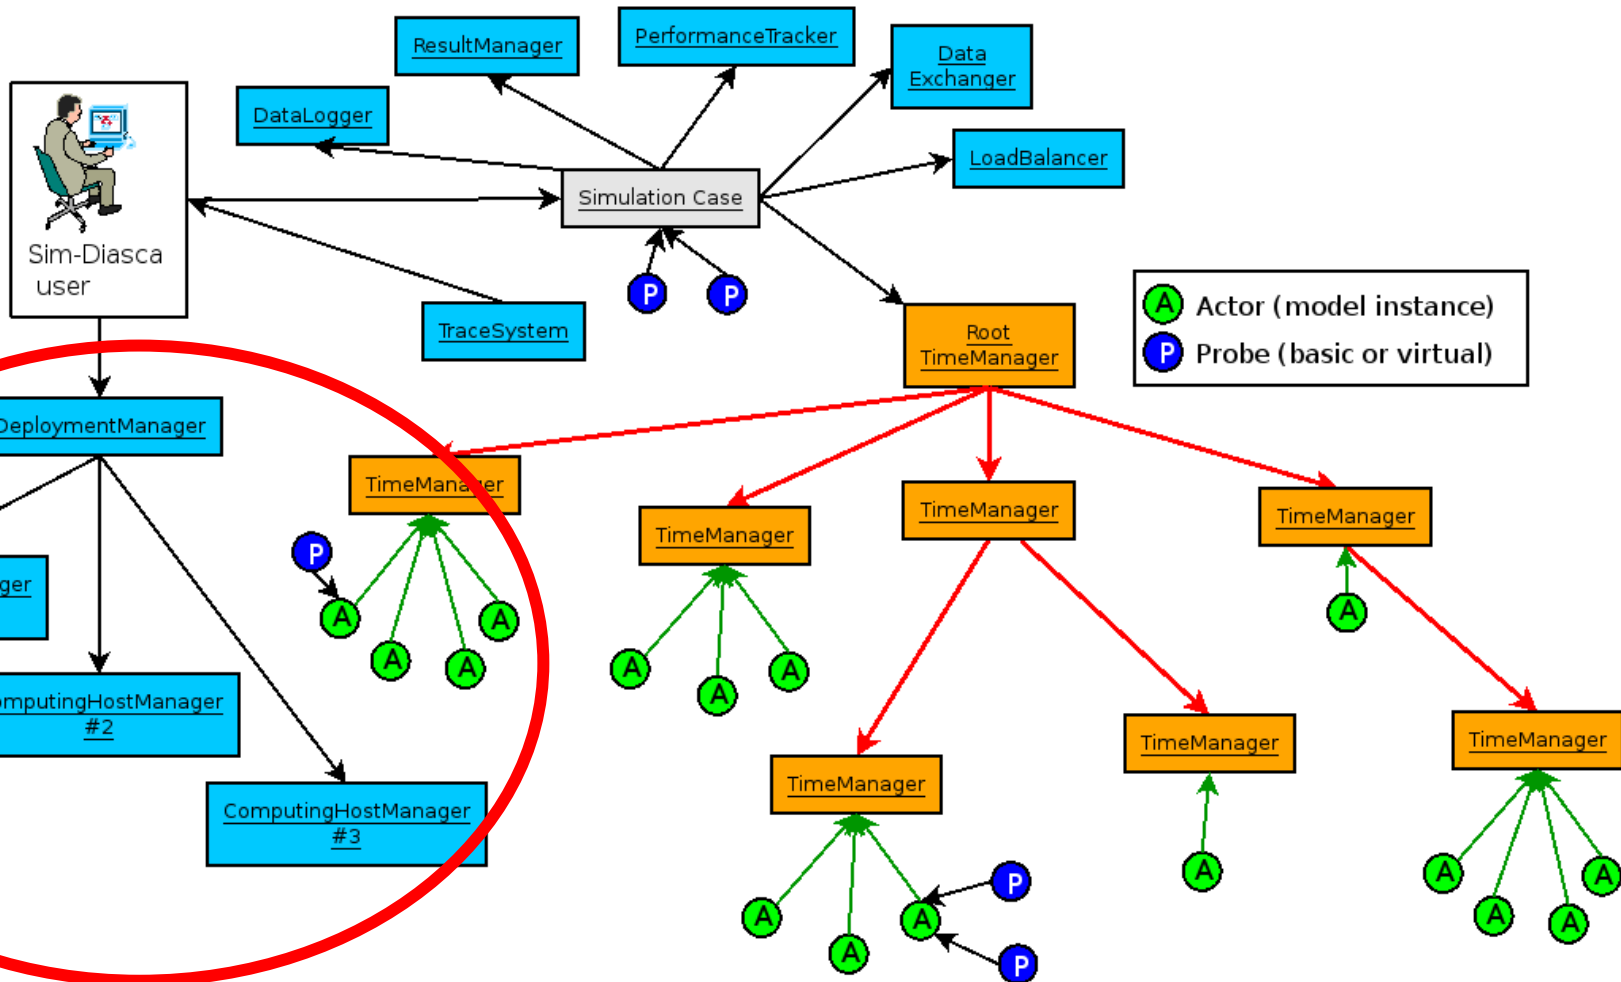
1: Read configuration information



Functional Services:

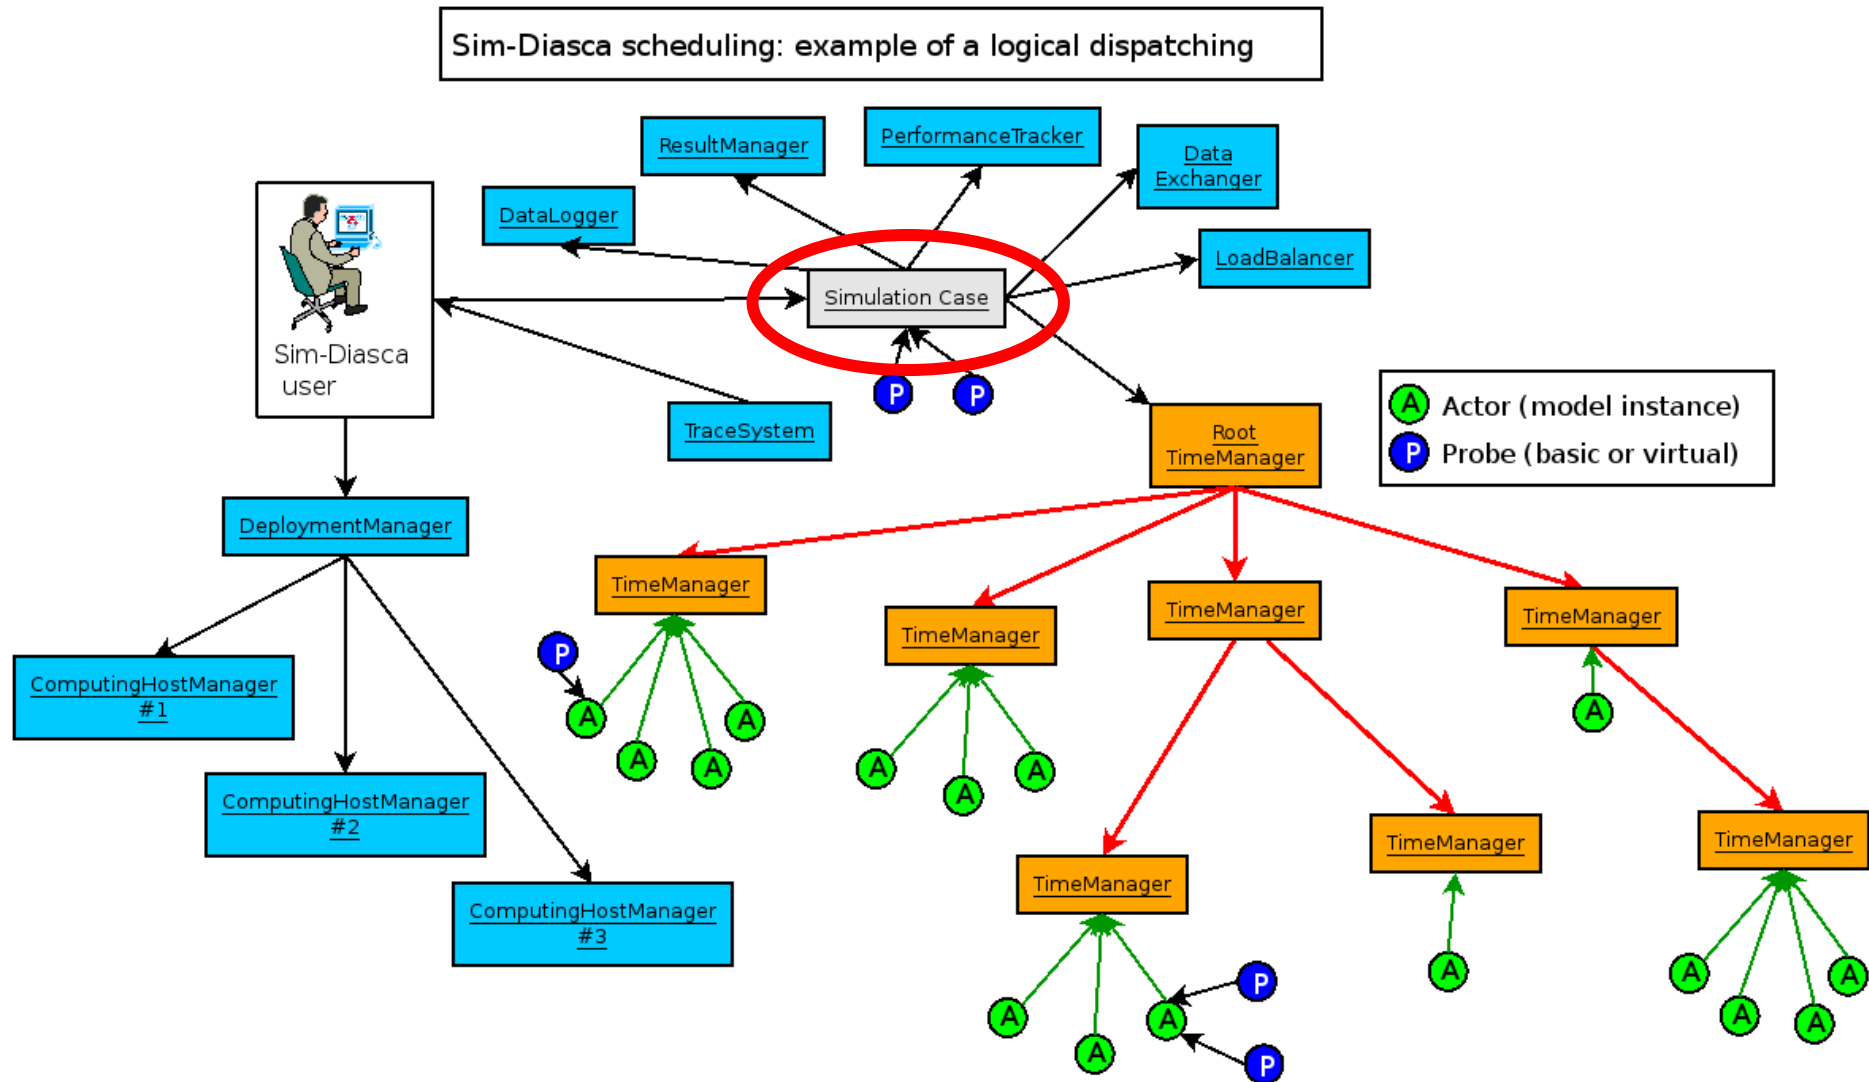
2: Deploy simulation on hosts
(no prior install needed)

Sim-Diasca scheduling: example of a logical dispatching



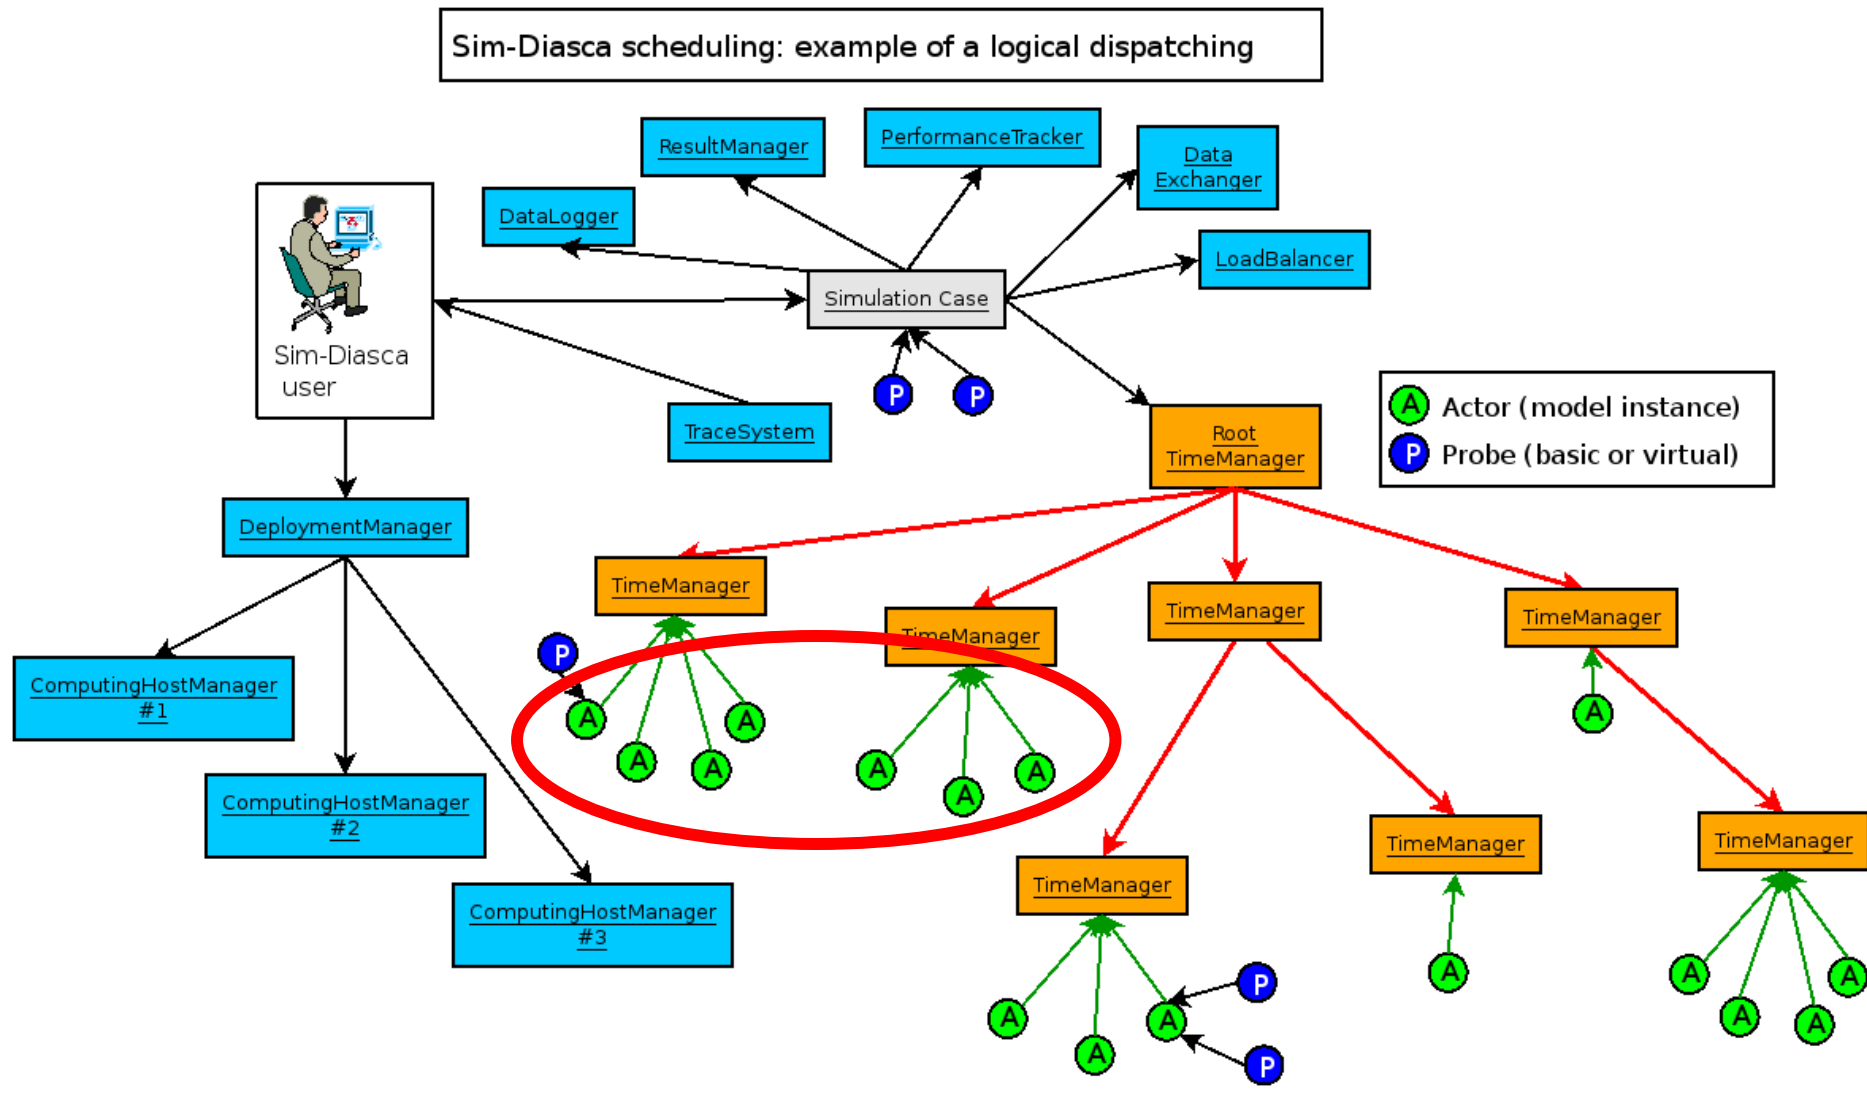
Functional Services:

3: Process simulation case



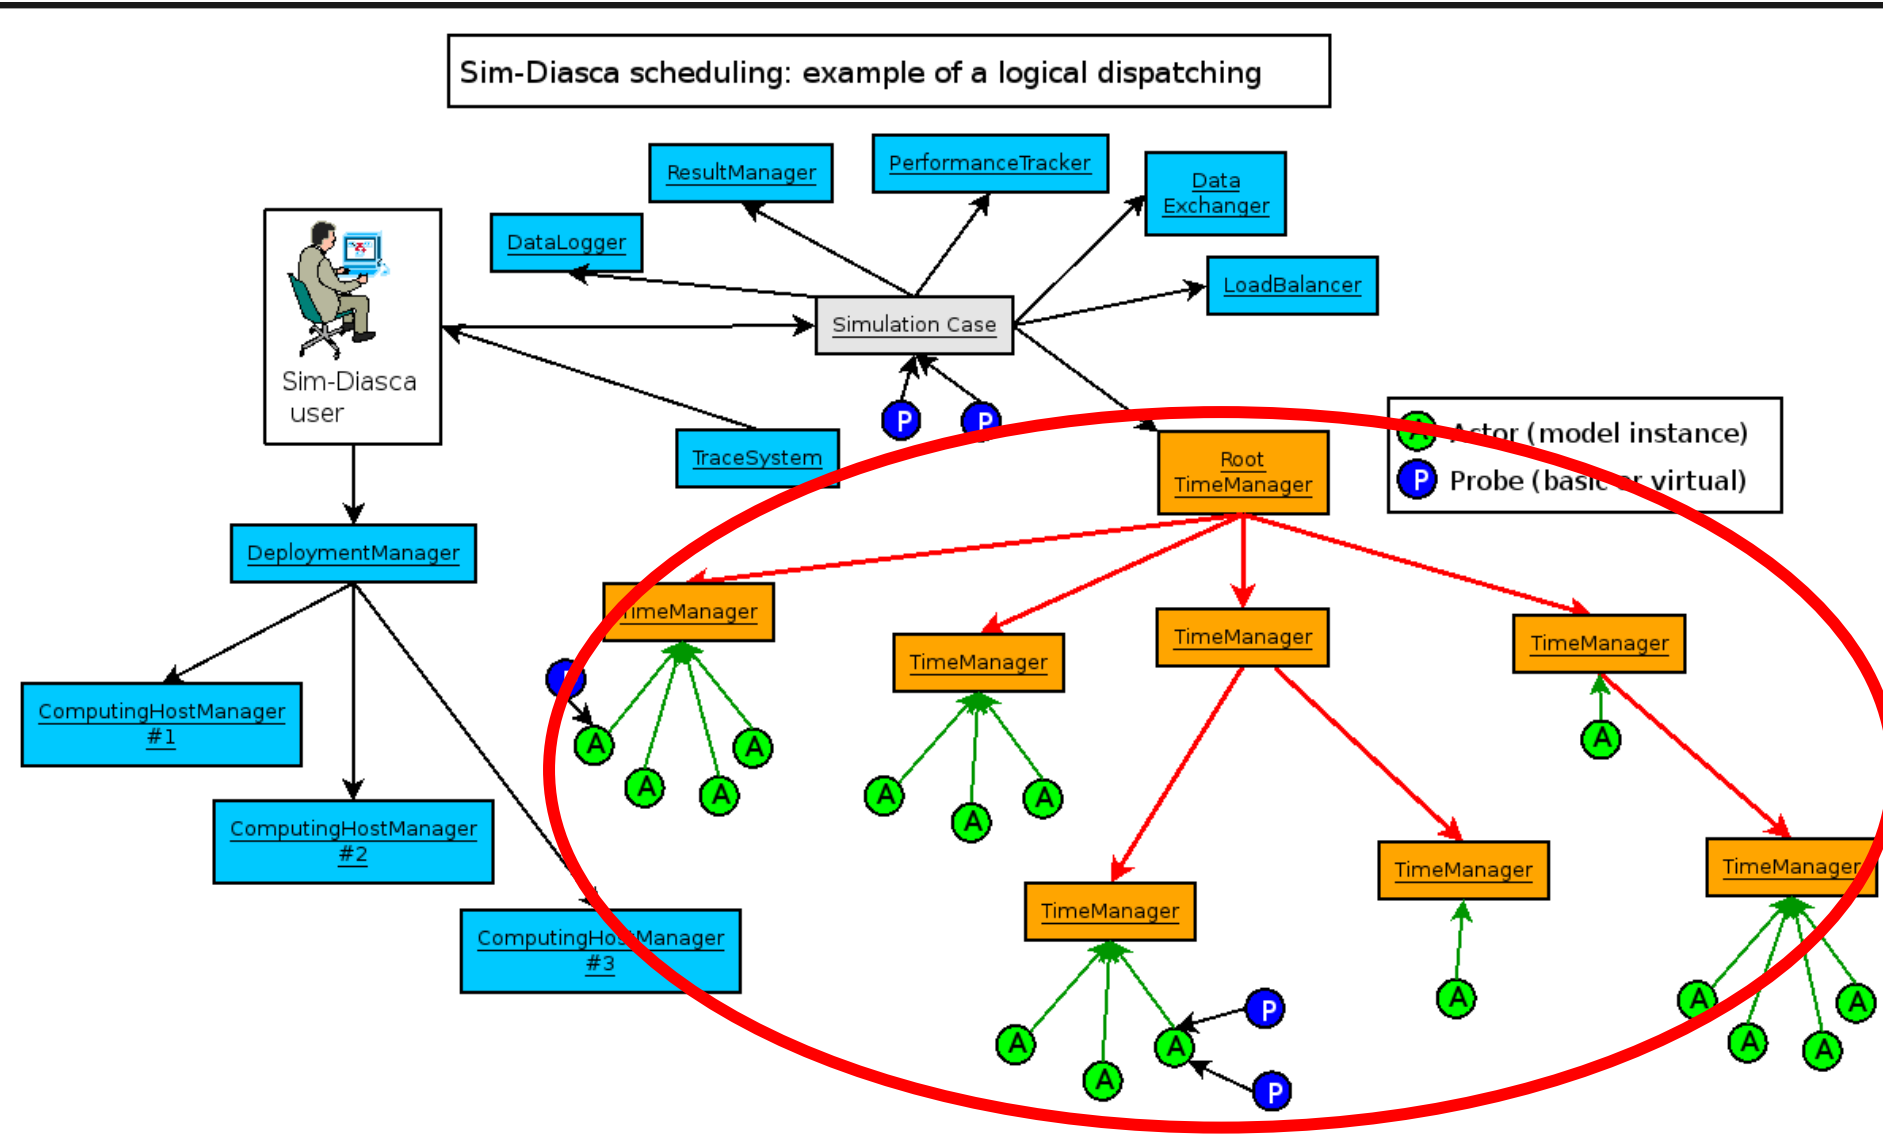
Functional Services:

4: Place and create initial model instances



Functional Services:

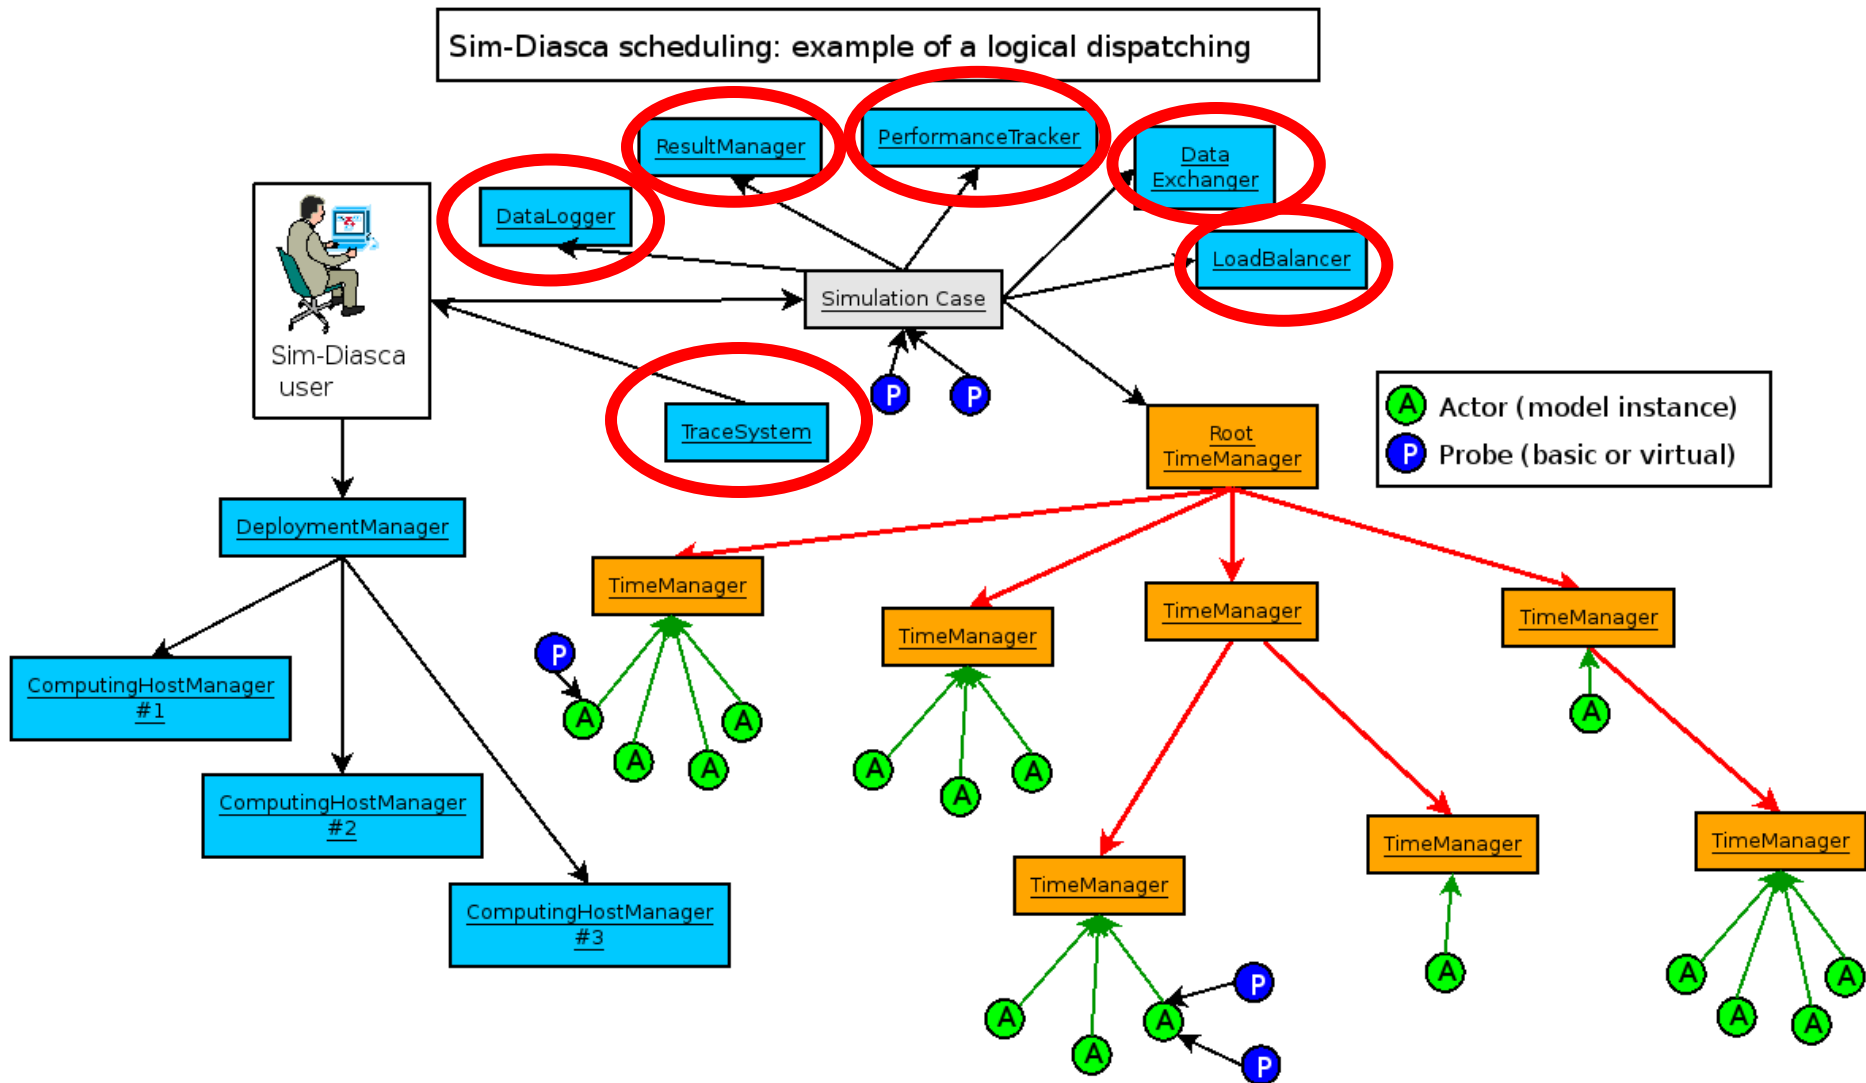
5: Schedule model instances



Functional Services:

6: Offer additional services

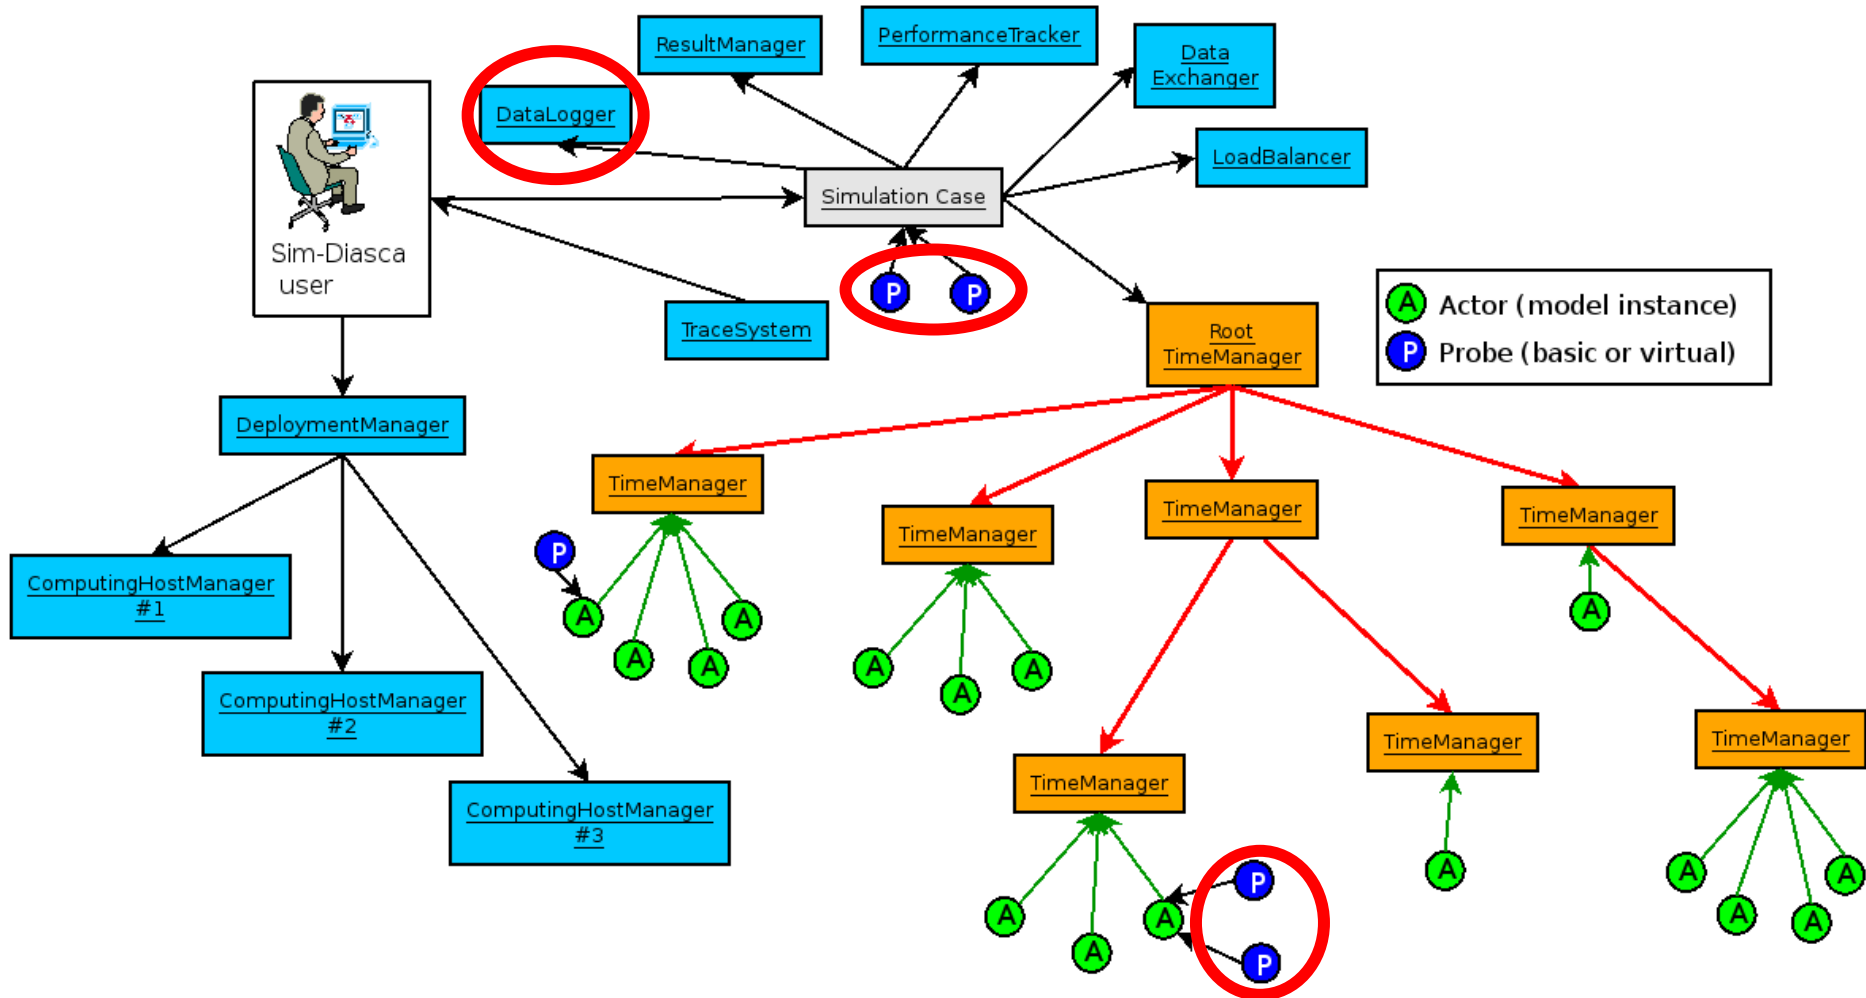
(e.g. distributed trace system, stochastic support, performance tracker, data-exchanger)



Functional Services:

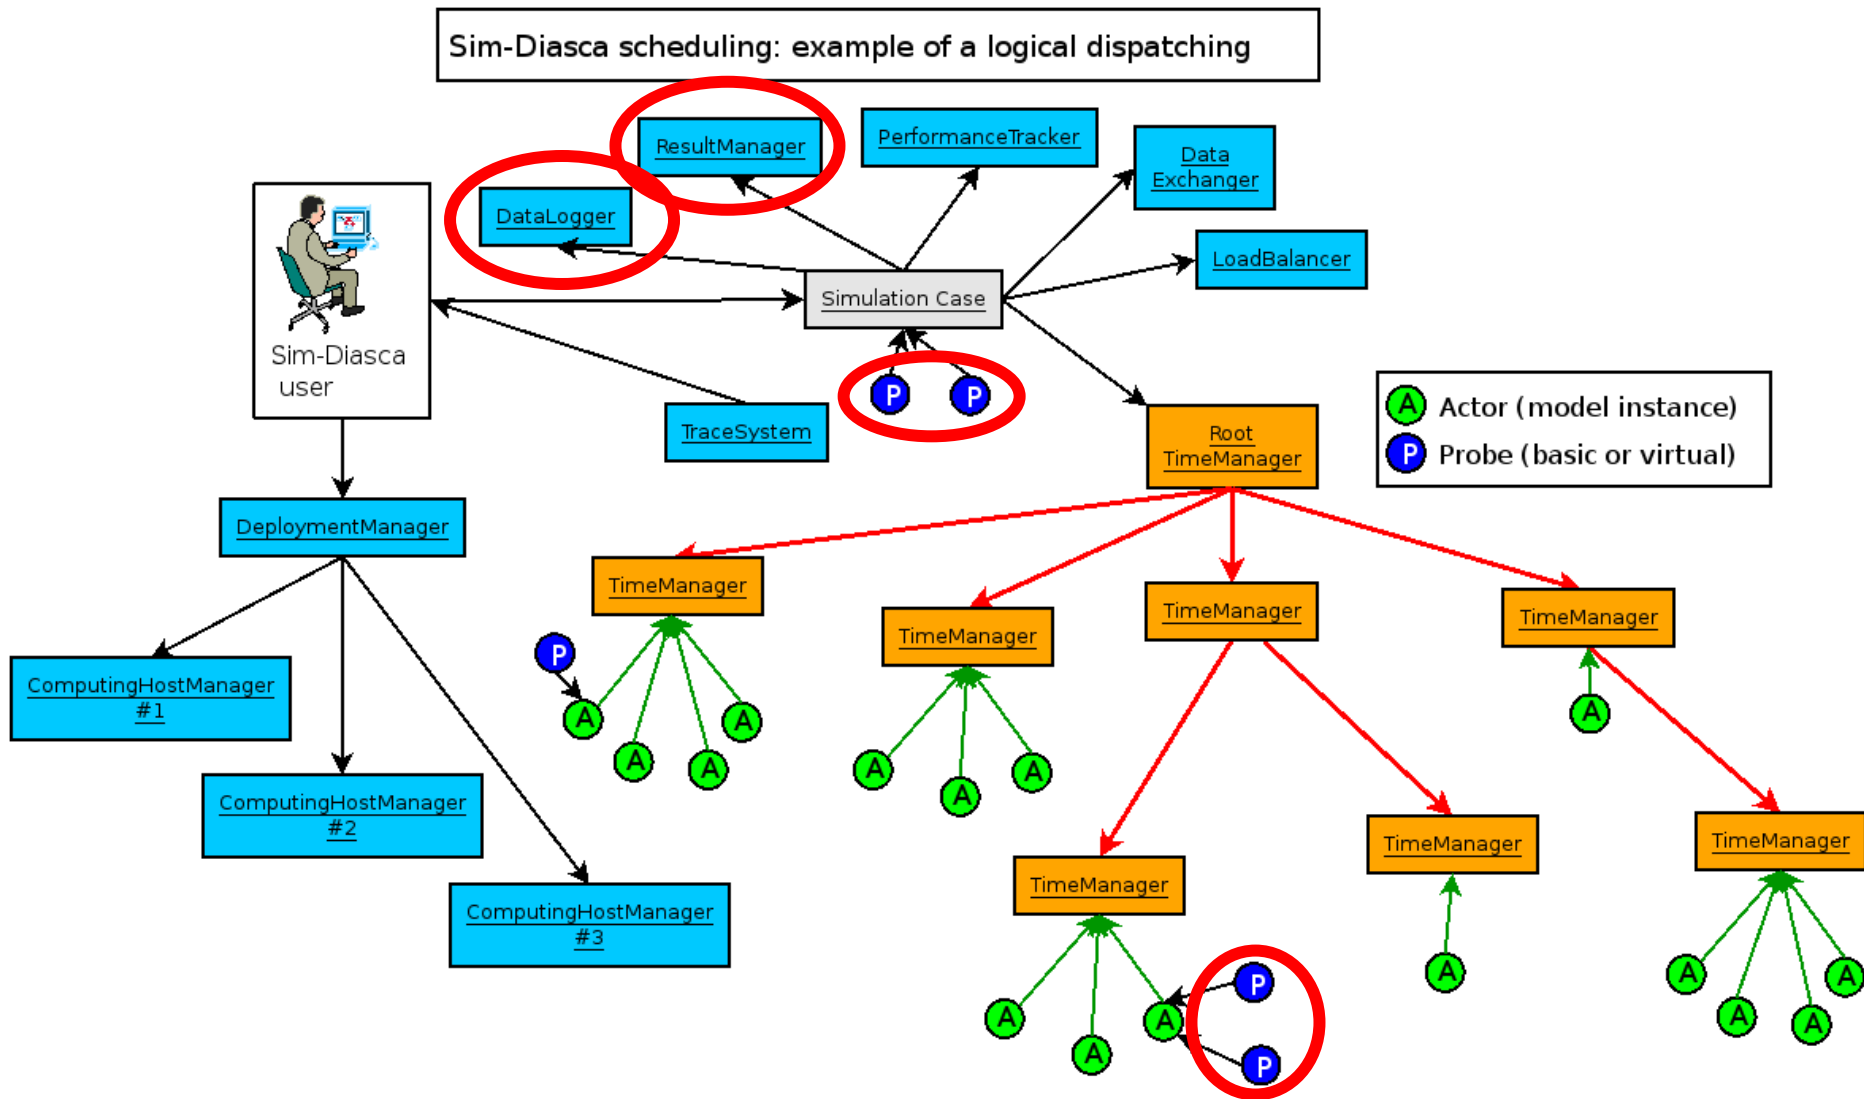
7: Generate simulation data (e.g. time series)

Sim-Diasca scheduling: example of a logical dispatching



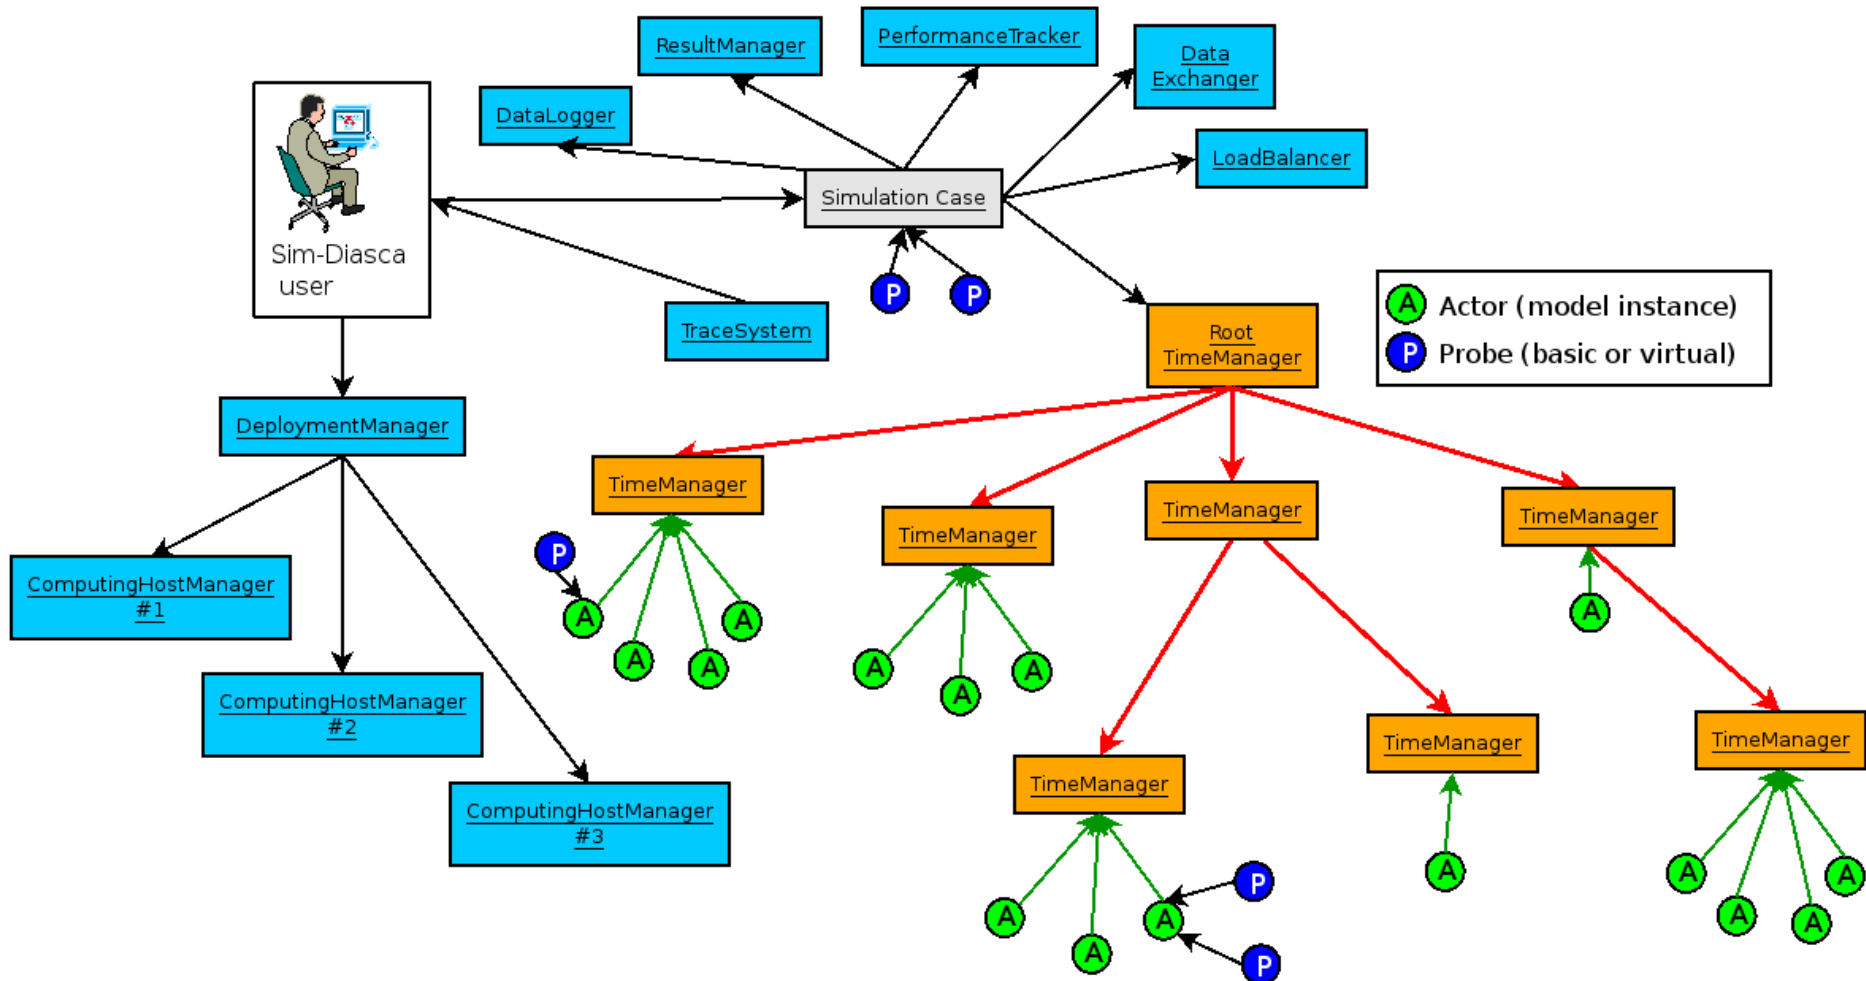
Functional Services:

8: Select, generate and retrieve results (e.g. plots)



Functional Services: (post-processing skipped)

Sim-Diasca scheduling: example of a logical dispatching



Sim-Diasca Key Points

- The engine must enforce notably following simulation properties:
 - Respect of causality
 - Total reproducibility
 - Some form of ergodicity
- The challenge is to obtain these properties in spite of massive parallelism and distribution
- To do so, inter-actor messages have to be appropriately reordered (transparently done by the engine)
- We target maximum scalability: potentially millions of instances of rather complex models can be evaluated in parallel
- Another technical challenge is to simplify as much as possible the model development:
 - **Bridge the gap** between model formalisations and actual simulation code (UML sequence diagrams, flow maps, state machines, dataflows): mixing programming styles (Object-Oriented and Functional ones)
 - **Provide stochastic support**
 - **Shelter models from parallelism**: write each of them in a simple, purely sequential way (and support a few bindings)
 - **Provide an higher-level language** with appropriate constructs so that domain experts have a better chance of understanding and developing actual models

More precise functional and technical requirements are detailed in annex 2.

Simulation properties discussed in annex 5.

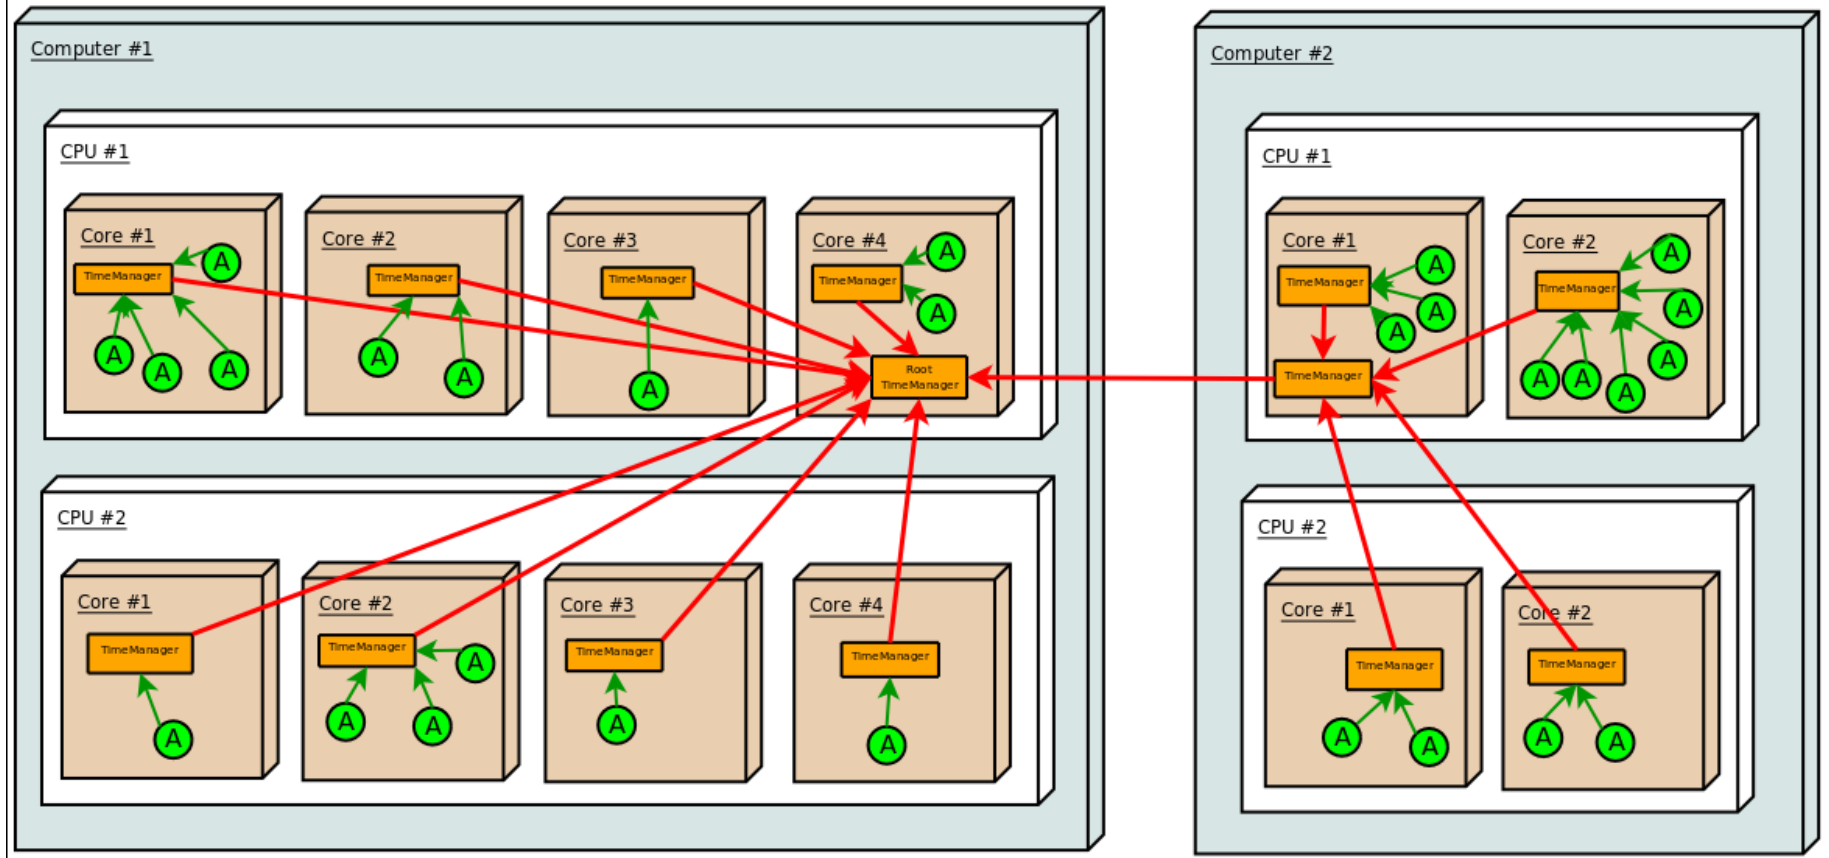
Simulation class and algorithmic aspects are better explained respectively in annex 3 and annex 6.

Overall modelling and simulation approach described in annex 7.

Refer to the *Sim-Diasca Technical manual* for further information.

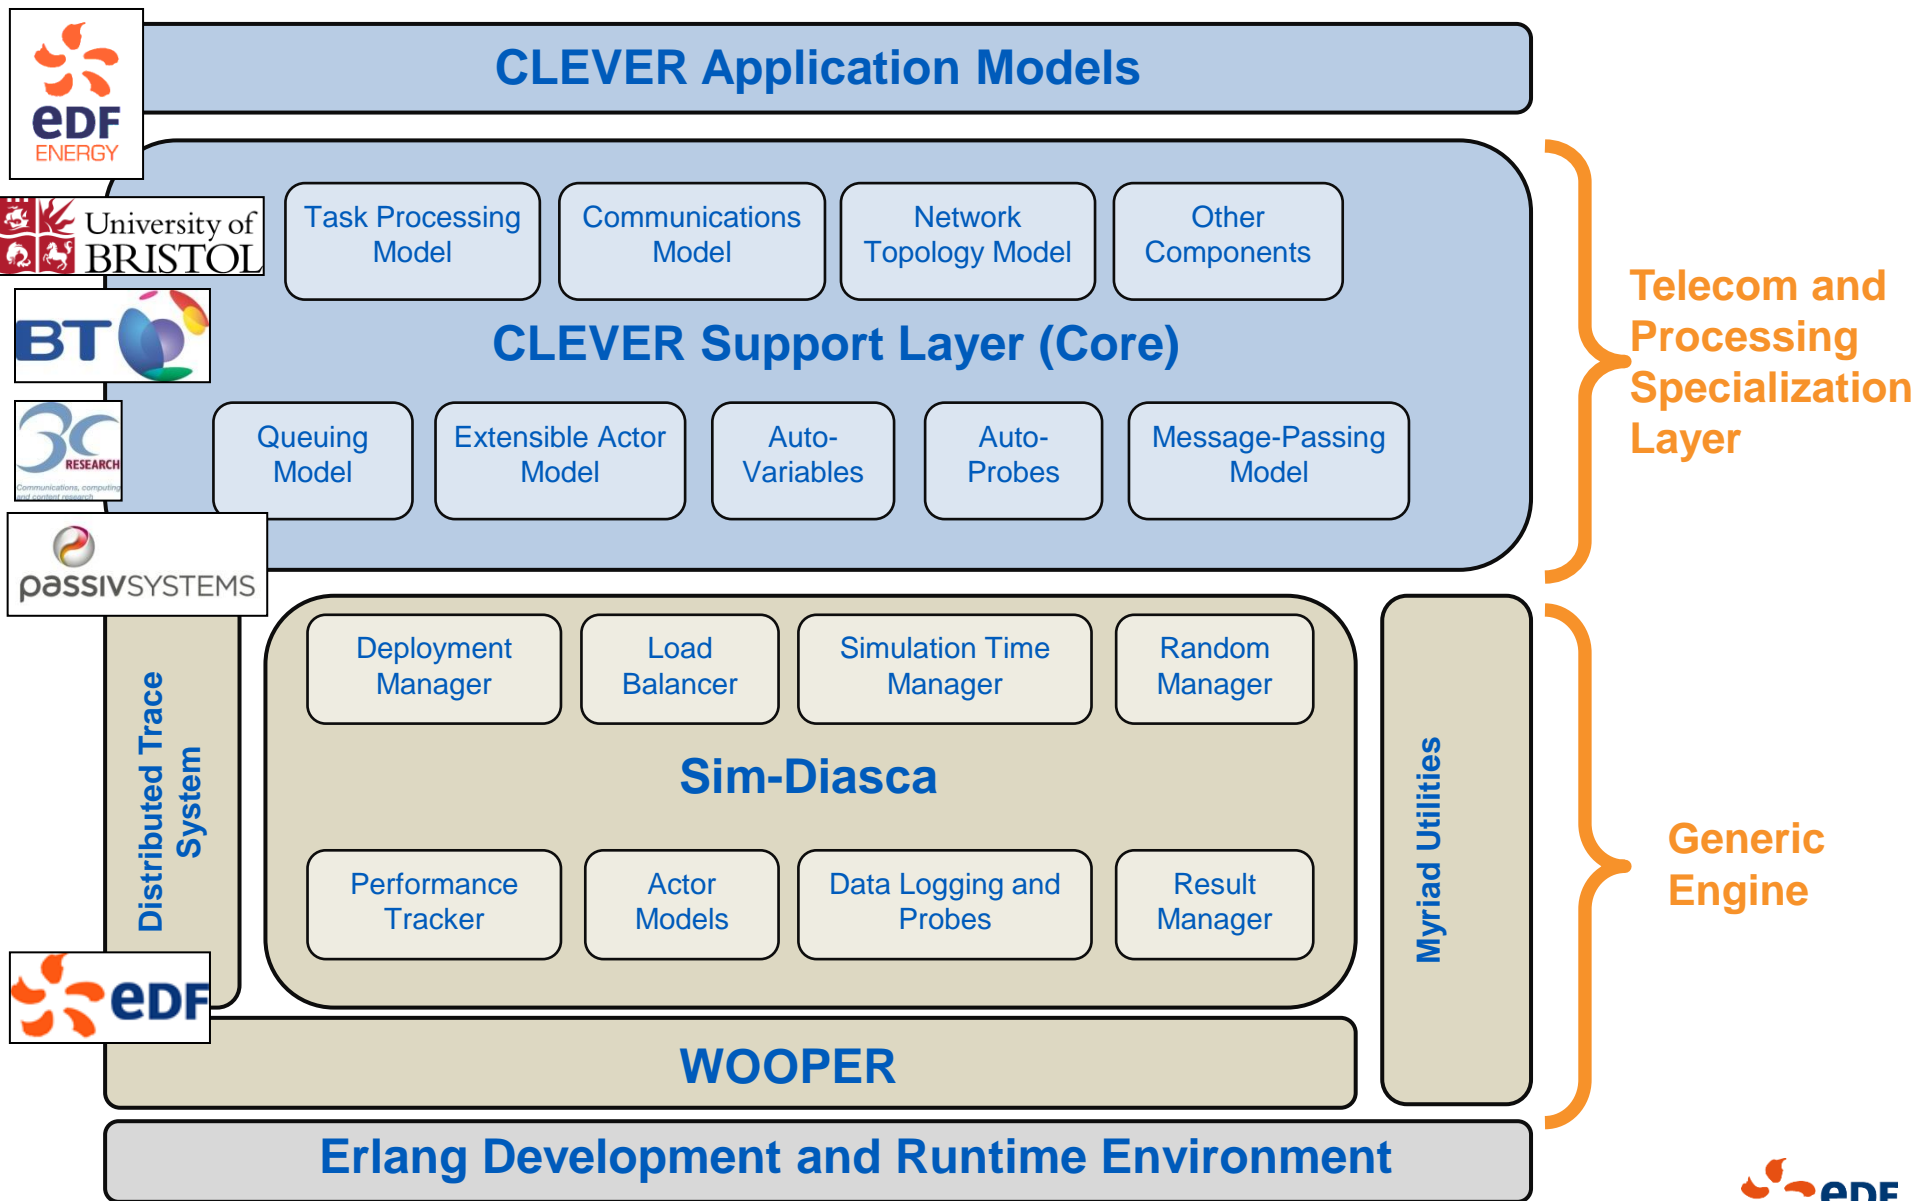
Sim-Diasca Technical Architecture in Practice

Sim-Diasca scheduling: example of a possible physical dispatching



- Most simulation services are at least partially distributed
- Dependencies on latency and bandwidth have been minimized (e.g. placement hint, advanced scheduling, hierarchical aggregation)
- High Performance Clusters supported, other platforms have been investigated (Tilera manycore cards, Bluegene/Q supercomputers)

Example of a technical architecture based on Sim-Diasca: the CLEVER simulator



Sim-Diasca: past, current and next steps

- After having been developed initially for the French case, used in the **CLEVER UK project**, in the **RELEASE european project**, in the **EDF City Platform** for the **MUG Project** (not counting the external, non-EDF uses):
- Possible follow-ups for:
 - The British supplier case and/or for Ofgem (regulator), i.e. a « CLEVER 2.0 »
 - The French counterpart supplier and/or DNO (distribution operator)
 - Smart grids (bridge towards equational models in continuous time, based on FMUs) and « future internet » related projects
 - Projects about scalability/reliability: some other Complex Systems of interest
(outside of the strict energy field, like urban planning, operational use of blockchains and DLT, intricate large-scale planning verification, digital twin of information systems)
- On the technical side:
 - Plenty of improvements could be considered (e.g. half-word emulator, more metaprogramming, hibernation, native compilation, Rust binding)
 - Larger-scale computing resources and k-crash resilient engine
- On the theoretical/academic side:
 - Scalability, reliability and performances to be investigated at this level too (e.g. with Coq?)
 - Towards hybrid simulations, mixing discrete time and continuous time with ODEs?
 - More generally: rising interest in functional programming for the scientific field
(e.g. EDF-CEA-INRIA 2012 Summer School)
 - Even if coming from an industrial background, some opportunities of publications

Ending word

- ▶ Sim-Diasca: a fully functional simulation engine, already used in academic works and industry-related projects
- ▶ Very few scalable engines of that kind exist (most are sequential by design)
- ▶ For a better understanding of how Sim-Diasca works, read the next appendices
- ▶ Fully generic: use it to simulate your own target system!
- ▶ Various paradigms supported: multi-agent simulation, dataflow evaluation; an additional one could be the support of at least some models in continuous time (hybrid mode of operation)
- ▶ Sim-Diasca has been released as free software (LGPL) by EDF R&D since 2010
- ▶ Main requirements are fulfilled, steady progresses expected to come

Towards a parallel, distributed, metaprogrammed Sim-Diasca running very large-scale hybrid simulations on larger HPC infrastructures?

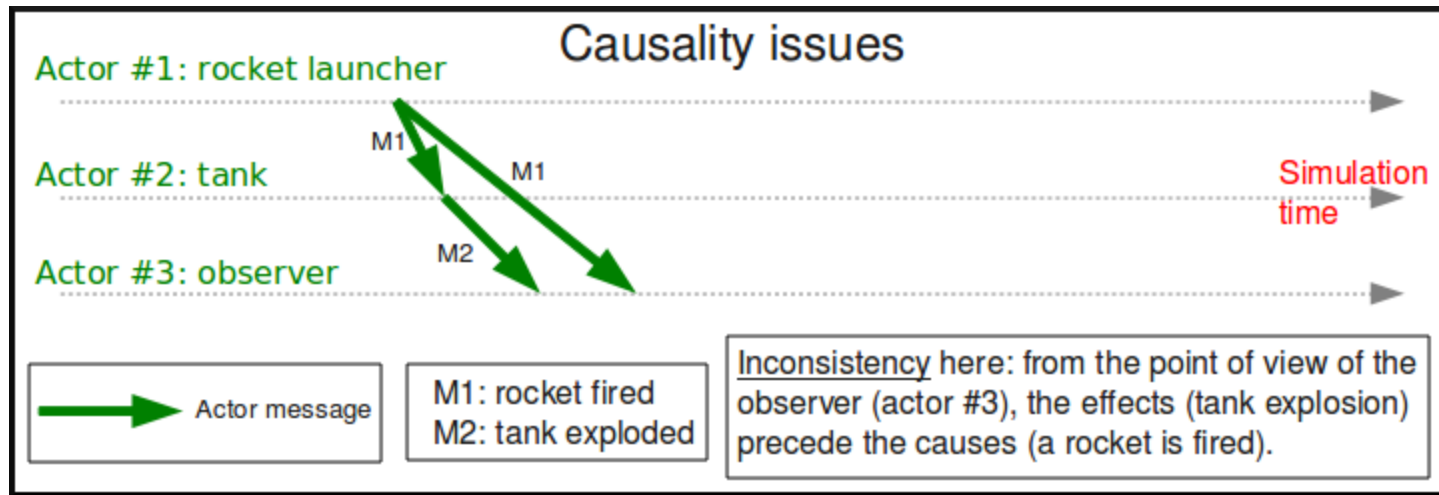


Appendices

- **Annex 1: Simulation properties**
 - 1.1: Preserving causality
 - 1.2: Ensuring total reproducibility
 - 1.3: Obtaining ergodicity
- **Annex 2: Mode of operation**
- **Annex 3: Anatomy of a Virtual Experiment**
- **Annex 4: Dataflow support**
- **Annex 5: Software-level considerations**
- **Annex 6: A Focus on WOOPER**
- **Annex 7: Overall modelling and simulation approach**
- **Annex 8: Examples of outputs**

Annex 1.1: Obtaining targeted simulation properties: restoring causality

By default, there is no total order of events over a distributed system (as no global time can exist). If no specific order is enforced, no consistency can be guaranteed:



Causality is natively respected in Sim-Diasca thanks to a time-stepped approach:

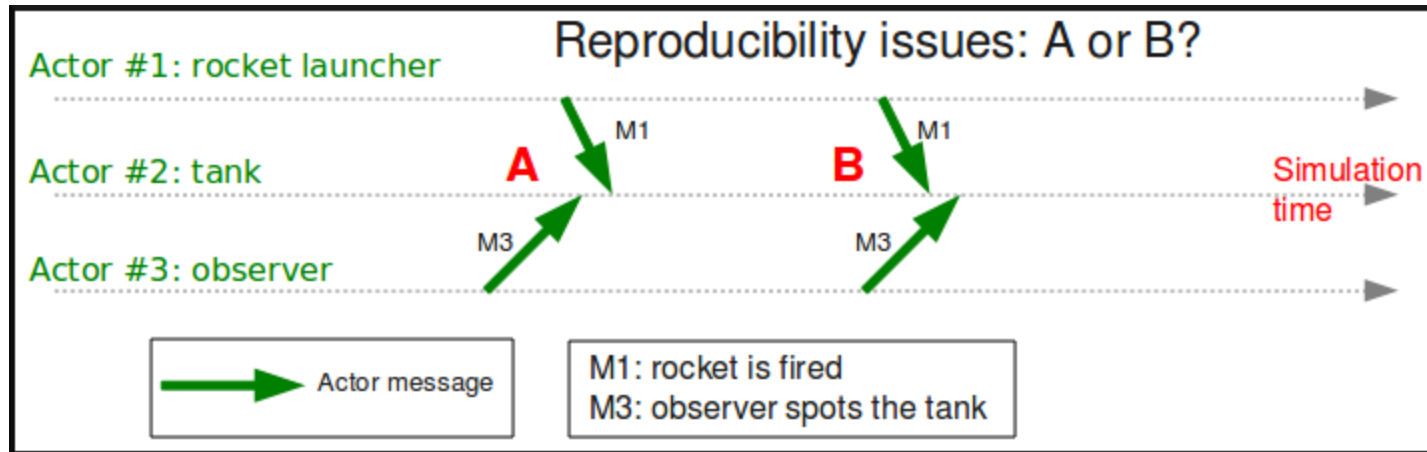
- ▶ To each cause corresponds necessarily an inter-actor message sent by actor A to actor B at tick T (*the physical time in the simulation*), diasca D (*a logical moment within a tick*)
- ▶ This message is processed (to determine its consequences) by actor B at tick T, diasca D+1

So, by design, as diascas do not overlap, causes indeed precede effects.

But causal chains induce only *partial ordering* of events.

What about *concurrent* events, i.e. events not linked by a causal relationship?

Annex 1.2: Obtaining targeted simulation properties: allowing for total reproducibility



No *a priori* order exists between two concurrent events.

(« M1 then M3 » is not any truer than « M3 then M1 »)

◆ The order that is to be re-created will necessarily be arbitrary.

Relying on the actual, technical order of receivings would make simulations depend on their execution context, and they would not be reproducible.

◆ The order that is to be re-created will have to fully abstract out any technical context.

With Sim-Diasca, each actor is reproducibly seeded and starts a diasca by automatically reordering the messages it received on the previous diasca, based only on:

- a reproducible identifier of the sending actor
- a hash value of the content of the message having been sent

This arbitrary reproducible order is generic and fully compatible with parallelism & stochastic models.

Annex 1.3: Obtaining targeted simulation properties: allowing for ergodicity

Objectives:

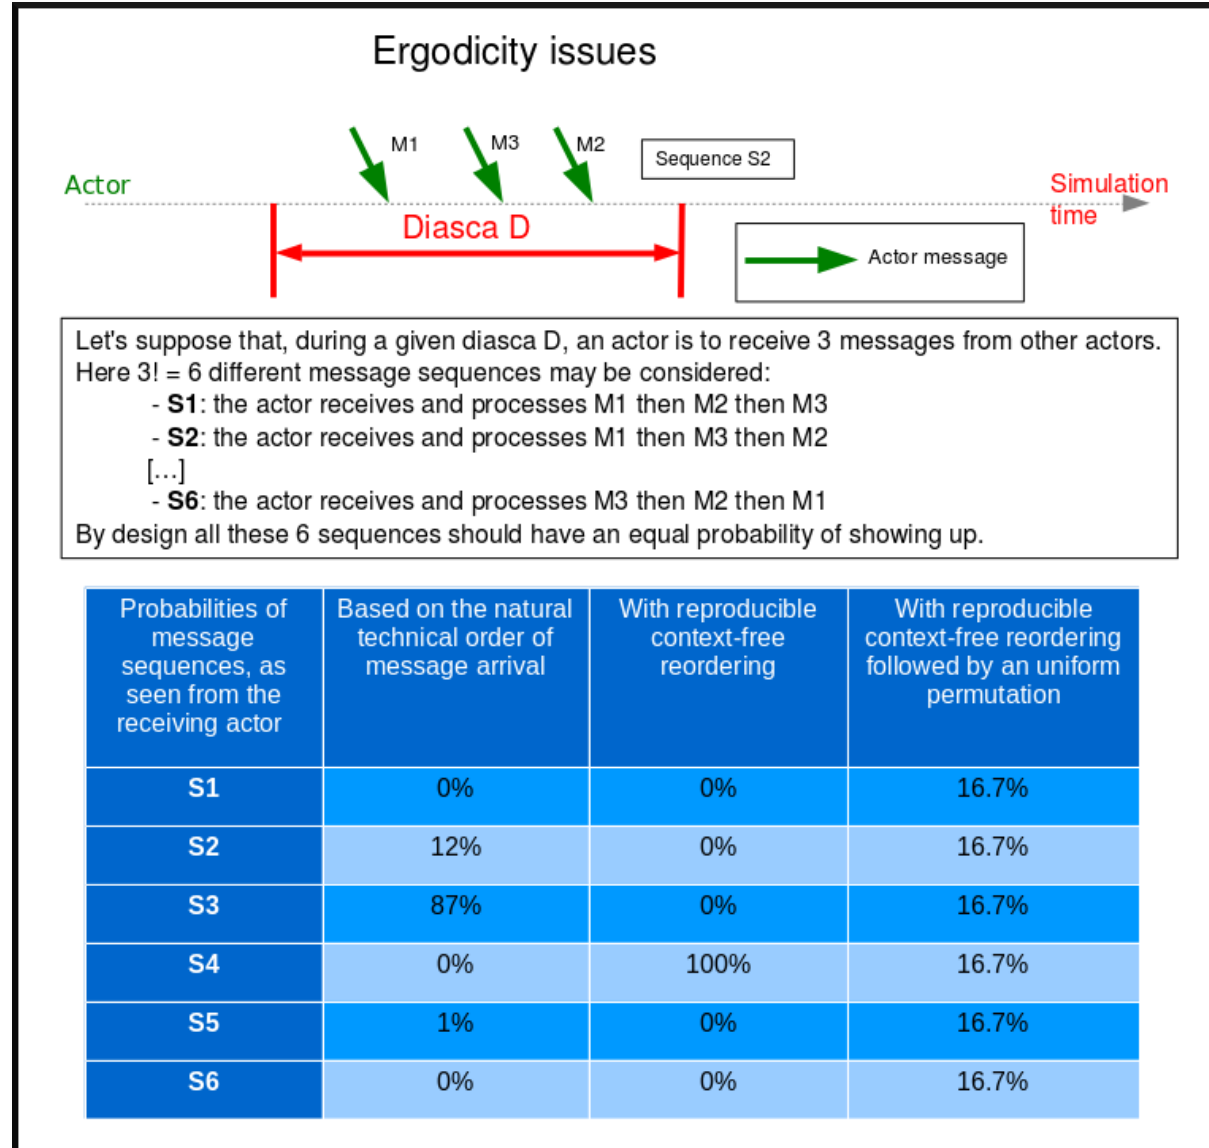
A. Ensure at each diasca, for each actor, that:

- for n messages received on the last diasca, all $n!$ possible sequences S_k can occur
- And that $P(S_k) = 1/n!$ for $k \in [1, n!]$

B. Ensure that running a given simulation case with a given root initial random seed will fully determine the simulation outcome, for a given set of parallel and/or distributed resources involved


Sim-Diasca approach, at the level of each actor, in parallel:

- The sequence of n past messages is first arbitrarily reordered for reproducibility, as shown previously
- Then, thanks to a distributed and uniform random generator, we select (reproducibly) one of the $n!$ possible permutations, and process it in-order
- Both objectives are then met: ergodicity (A) and reproducibility (B)



Annex 2: Sim-Diasca mode of operation

Main objective: massive parallelism while still preserving the expected simulation properties. Should you have 35 million model instances, you are here able to evaluate all of them *in parallel*.



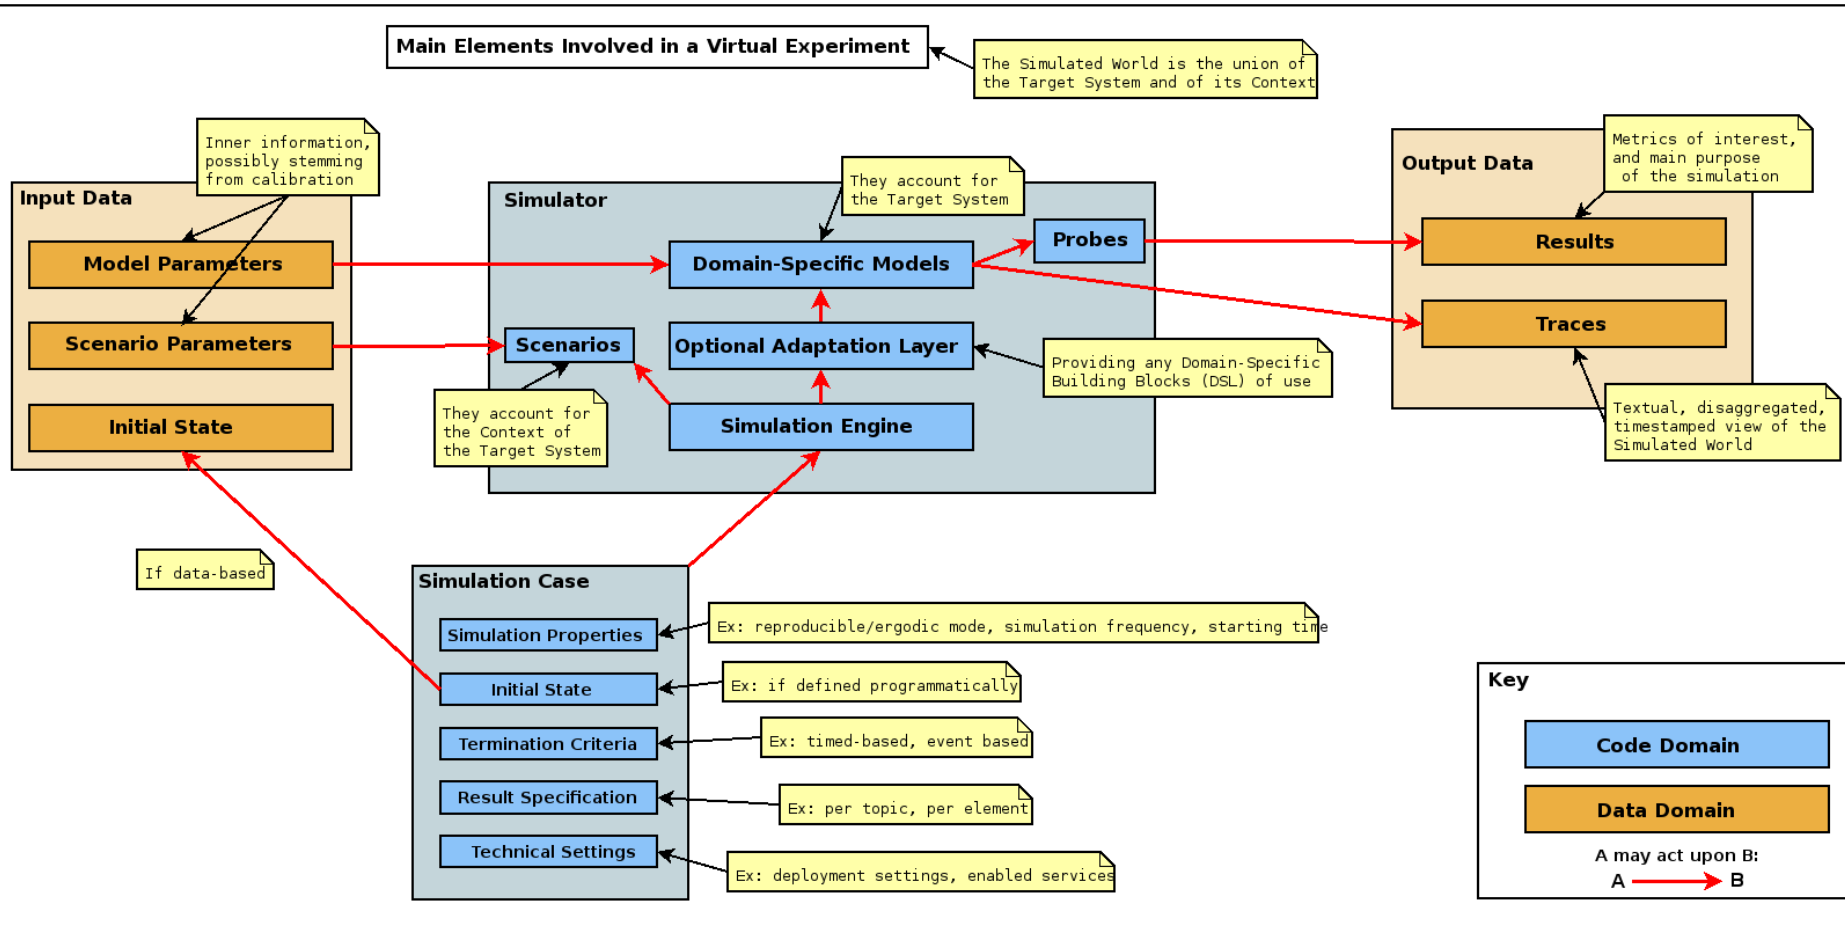
Note: this diagram shows how Sim-Diasca, whose simulation time is a (Tick, Diasca) pair, would evaluate the causality example in annex 1.1.

Following topics are not illustrated here:

- distributed mode of operation with hierarchical time managers
- actor-level reordering of messages
- actual life-cycle management
- stochastic management

- ◆ Simulation ticks are evaluated sequentially (one after the other, based on a uniform synchronous simulation time; yet only the necessary ones), each of their diascas evaluating its actors fully in parallel
- ◆ At diasca $D=0$, all actors (model instances) having planned a spontaneous behaviour execute it; at $D>0$, those having received inter-actor message(s) at $D-1$ reorder and process them
- ◆ Besides parallelism, two main goals to be achieved transparently by the engine:
 - To obtain a consensus on the correct soonest termination of each tick and diasca, thanks to a necessary and sufficient exchange of synchronisation messages (involving actors and time managers)
 - To reorder automatically received messages so that the targeted properties (causality, reproducibility, ergodicity, etc.) are preserved

Annex 3: Anatomy of a Virtual Experiment

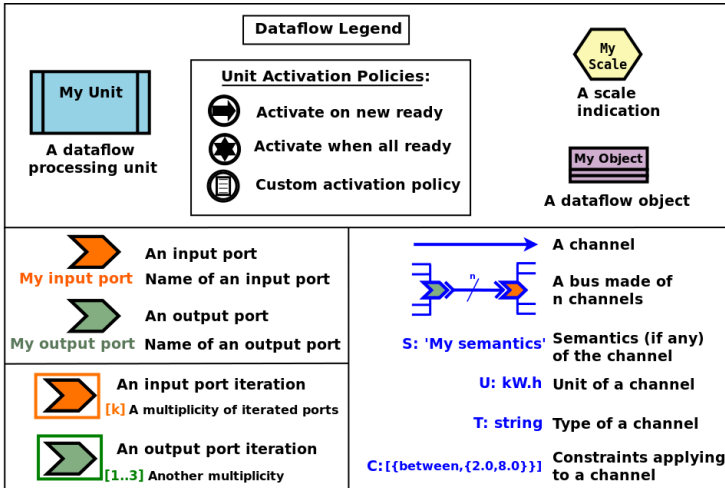


(refer to our mini-ontology for more detailed descriptions, and to `mock-simulators/soda-test/src` for a runnable example thereof)

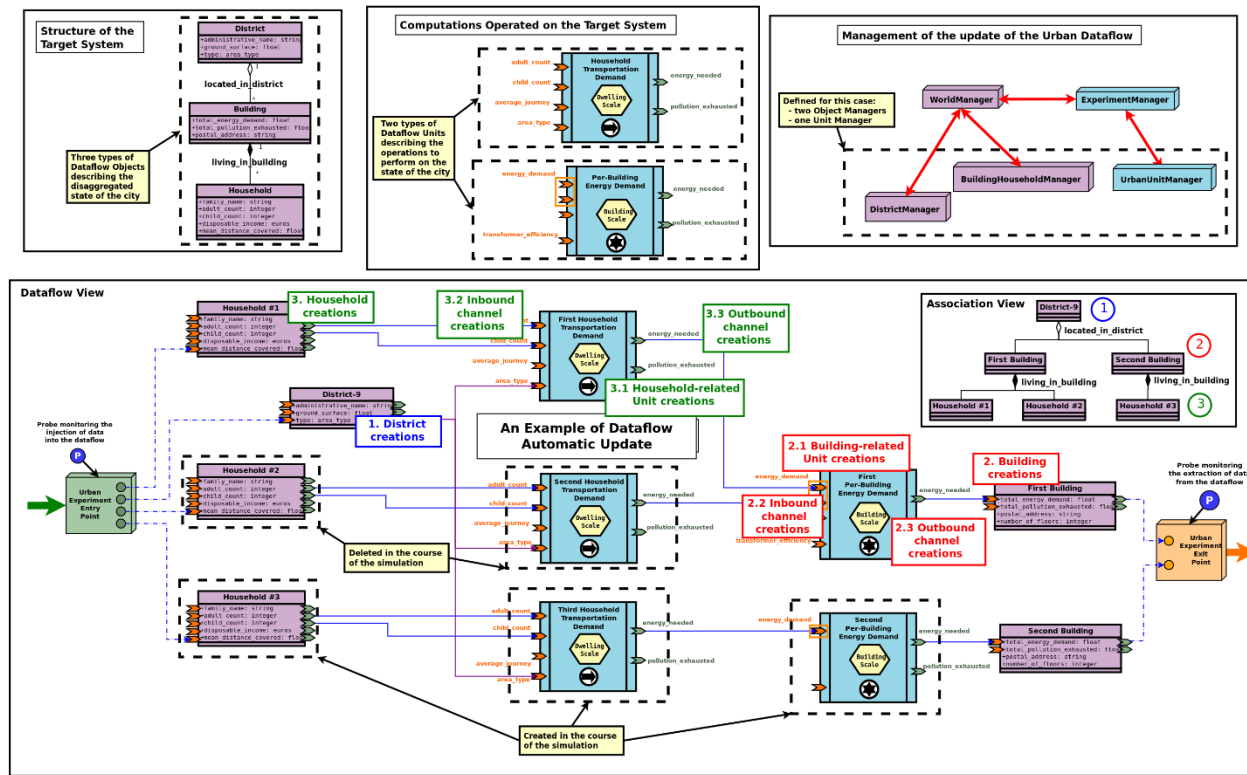
Annex 4: Dataflow support (1/2)

Sim-Diasca natively powers multi-agent simulations.

Among the specialisations that can be built on it, one deals with the parallel execution of dataflows, i.e. *graphs of computations whose evaluation is driven by the availability of inputs.*





Dataflow corresponding to the Urban-Example simulation case



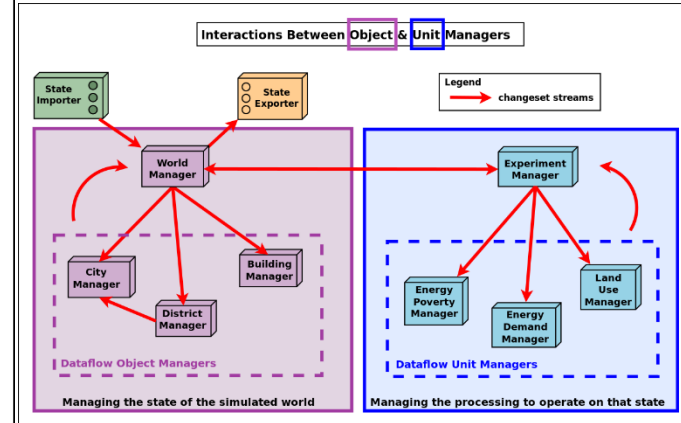


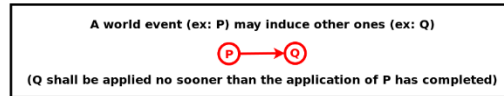
Annex 4: Dataflow support (2/2)

These dataflows can be cyclic and highly dynamic: during a time-step, any number of blocks, ports and channels can be created or destroyed, based on object managers and unit managers, which exchange changesets:

8 types of world events may affect the blocks of a dataflow



Describing and applying dataflow changes using world events



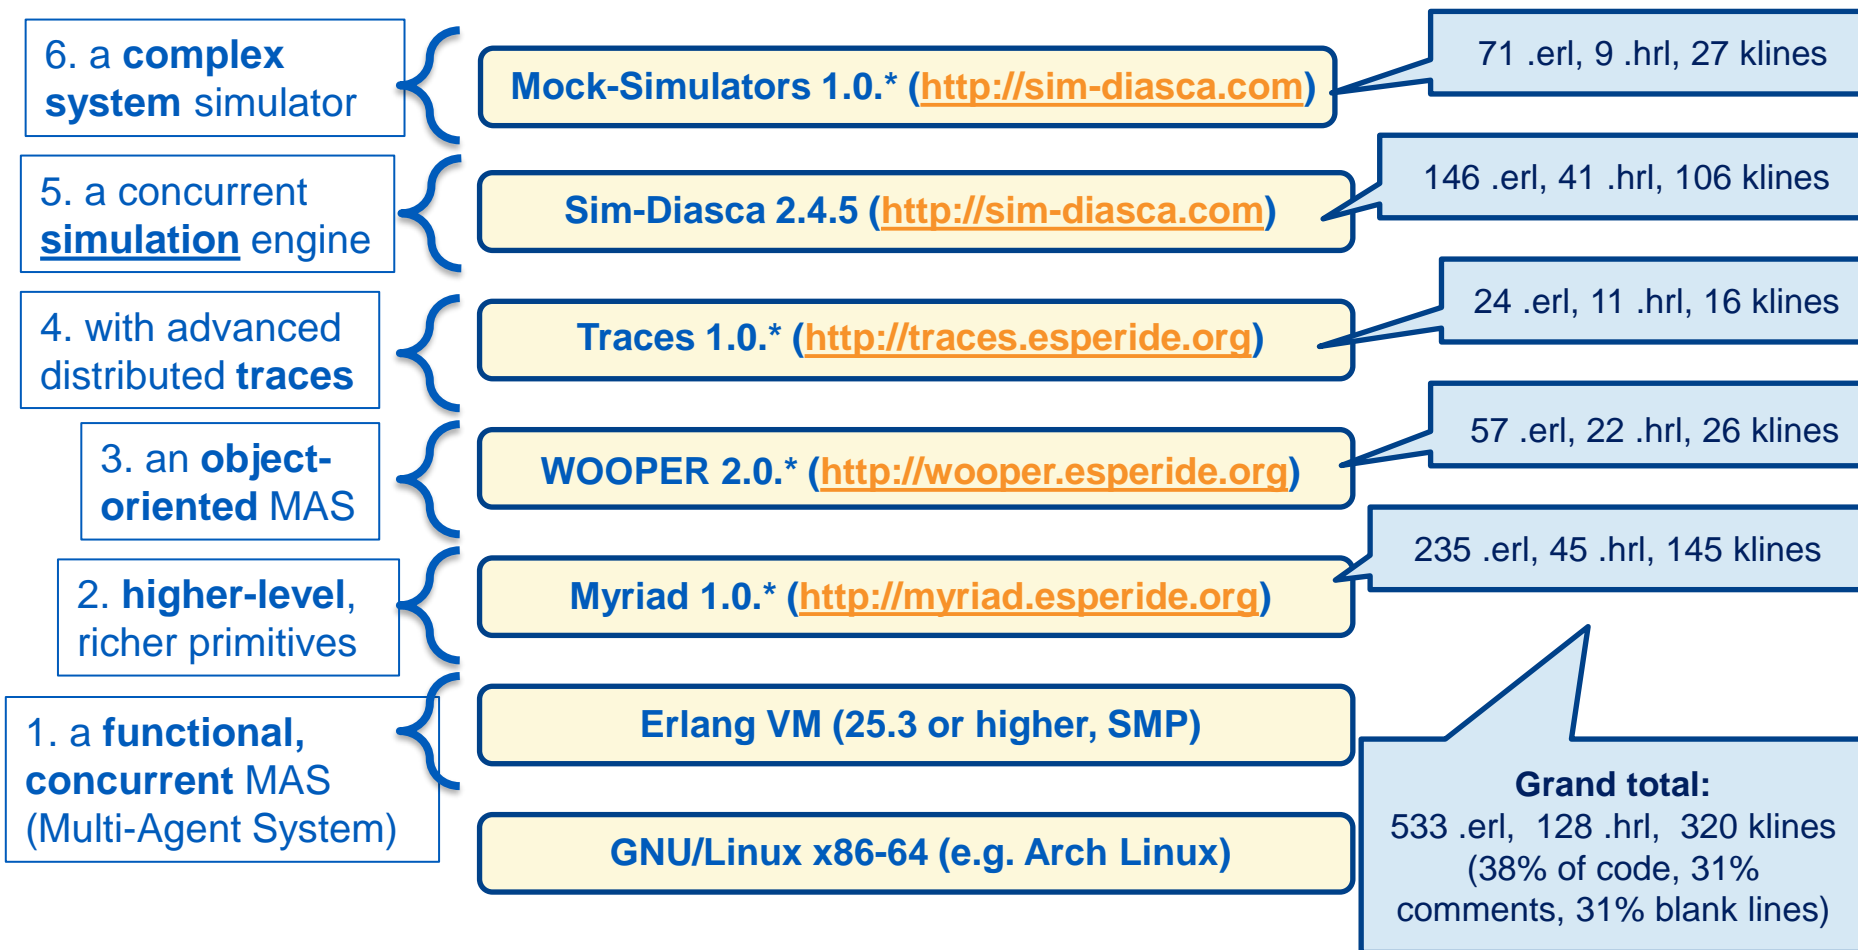
Language bindings have been defined (in Python; Java considered), so that processing units may be developed in other languages and/or embed third-party pre-existing models.

A workbench has been designed to simplify such integrations.

Annex 5: Software-level considerations

A more technical zoom on the open source software stack

Through a bottom-up specialisation, obtaining in turn:



Annex 6: A Focus on WOOPER (1/4): Wrapper for Object-Oriented Programming in Erlang

- ▶ A free-software lightweight OOP layer on top of the [Erlang](#) language
(now OTP-compliant thanks to rebar3, for applications/releases, and available as an Hex package)
- ▶ Official website: <http://wooper.esperide.org>
(sole dependency: Myriad, see <http://myriad.esperide.org>)
- ▶ Provides multiple inheritance, polymorphism, encapsulation, state management and life-cycle management with minimal development/runtime overhead
- ▶ You can define classes, create instances, call methods (oneway or requests) on them, and delete them
- ▶ A class definition includes: a list of the direct superclasses ([`class_Mammal`, `class_Viviparous`]), at least one constructor, possibly a destructor, and member and static methods (e.g. `declareBirthday/1`)

Thus provides an object-oriented, distributed multi-agent system; but such features alone are *not* sufficient to obtain a simulation!

Annex 6: A Focus on WOOPER (2/4): implementing a class

```
9
10 % @doc Cat-based example. Those are the ones that work best.
11 %
12 % Guaranteed to be implemented by a cat.
13 %
14 module(class_Cat).
15
16
17 -define( class_description,
18         "Class modelling any kind of cat, and there are many." ).
19
20
21 % Determines what are the direct mother classes of this class (if any):
22 -define( superclasses, [ class_Mammal, class_ViviparousBeing ] ).
23
24
25 -define( class_attributes, [
26     { whisker_color, whisker_color(), none,
27       "50 shades of whiskers" } ] ).
28
29
30 % Allows to define WOOPER base variables and methods for that class:
31 -include("wooper.hrl").
32
33
34 % Import common types without module prefix:
35 -include("ecosystem_types.hrl").
36
37
38 % Shorthands:
39
40 -type ustring() :: text_utils:ustring().
41
42
43 % @doc Constructs a cat instance.
44 -spec construct( wooper:state(), age(), gender(), fur_color(),
45                whisker_color() ) -> wooper:state().
46 construct( State, Age, Gender, FurColor, WhiskerColor ) ->
47
48 % First the direct mother classes:
49 MammalState = class_Mammal:construct( State, Age, Gender, FurColor ),
50 ViviparousMammalState = class_ViviparousBeing:construct( MammalState ),
51
52 % Then the class-specific attributes:
53 setAttribute( ViviparousMammalState, whisker_color, WhiskerColor ).
54
55
56
57 -spec destruct( wooper:state() ) -> wooper:state().
58 destruct( State ) ->
59
60 io:format( "Deleting cat ~w! (overridden destructor)~n", [ self() ] ),
61
62 State.
63
64
```

```
67
68 % @doc No guarantee on biological fidelity.
69 -spec getTeatCount( wooper:state() ) -> const_request_return( teat_count() ).
70 getTeatCount( State ) ->
71     wooper:const_return_result( 6 ).
72
73
74 % @doc Cats are supposed carnivorous though.
75 -spec canEat( wooper:state(), food() ) -> const_request_return( boolean() ).
76 canEat( State, soup ) ->
77     wooper:const_return_result( true );
78
79 canEat( State, chocolate ) ->
80     throw( { harmful_food_detected, chocolate } );
81
82 canEat( State, croquette ) ->
83     wooper:const_return_result( true );
84
85 canEat( State, meat ) ->
86     wooper:const_return_result( true );
87
88 canEat( State, OtherFood ) ->
89     wooper:const_return_result( false ).
90
91
92 % @doc Returns the whisker color of this cat.
93 -spec getWhiskerColor( wooper:state() ) -> const_request_return( color() ).
94 getWhiskerColor( State ) ->
95     io:format( "getWhiskerColor/1 request called by ~w.~n", [ ?getSender() ] ),
96     wooper:const_return_result( ?getattr(whisker_color) ).
97
98
99 % @doc Requests this cat to terminate, based on specified halting procedure.
100 -spec terminate( wooper:state(), 'crash' ) -> const_oneway_return().
101 terminate( State, crash ) ->
102     basic_utils:crash(),
103     wooper:const_return().
104
105
106
107 -spec toString( wooper:state() ) -> const_request_return( ustring() ).
108 toString( State ) ->
109     Description = text_utils:format( "cat instance with whiskers of color ~p.",
110                                     [ ?getattr(whisker_color) ] ),
111     wooper:const_return_result( Description ).
112
113
114
115 % Static section.
116
117
118 % @doc Returns the mean life expectancy of a cat, in years.
119 -spec get_mean_life_expectancy() -> static_return( age() ).
120 get_mean_life_expectancy() ->
121     wooper:return_static( 18 ).
```

Annex 6: A Focus on WOOPER (3/4): interacting with instances

```
MyC = class_Cat:new_link( _Age=3, _Gender=female, _FurColor=sand, _WhiskerColor=white ),
MyC ! { getClassName, [], self() },
receive
  { wooper_result, class_Cat } -> ok
end,
MyC ! { getSuperclasses, [], self() },
receive
  { wooper_result, _Classes=[ class_Mammal, class_ViviparousBeing ] } -> ok
end,
MyC ! { getAge, [], self() },
receive
  { wooper_result, 3 } -> ok
end,
MyC ! { setAge, 5 },
MyC ! { getAge, [], self() },
receive
  { wooper_result, 5 } -> ok
end,
MyC ! declareBirthday,
MyC ! { getAge, [], self() },
receive
  { wooper_result, 6 } -> ok
end,
MyC ! { canEat, soup, self() },
receive
  { wooper_result, true } -> ok
end,
MyC ! delete,
18 = class_Cat:get_mean_life_expectancy().
```


Annex 6: A Focus on WOOPER (4/4): Inner workings & Erlang mapping

- ◆ A WOOPER class is an (Erlang) module (e.g. `class_Cat`), a WOOPER (active) instance is an (Erlang) process, an instance identifier is a PID, method calls are messages, a state is a set of attributes, an attribute is a key/value pair (`{atom(), term() }`)
- ◆ A WOOPER instance is:
 - created thanks to `(timed_)(remote_)(synchronous_)new(_link)` calls
 - a process looping over its state, waiting for incoming method calls, mapping them to the corresponding functions in appropriate modules (inheritance), returning possibly a result
- ◆ It keeps a private associative table to hold its state (`attribute_name -> value`)

| WOOPER concept | Corresponding Erlang mapping |
|-----------------------|--|
| class definition | module |
| instance | process |
| instance reference | process identifier (PID) |
| new operators | WOOPER-provided functions, making use of user-defined <code>construct/N</code> functions (a.k.a. the constructors) |
| delete operators | WOOPER-provided functions, unless user-specified (a.k.a. the destructor) |
| method definition | module function that respects some conventions |
| method invocation | sending of an appropriate inter-process message |
| method look-up | class-specific virtual table taking into account inheritance transparently |
| instance state | instance-specific datastructure, kept by the instance-specific WOOPER tail-recursive infinite loop |
| instance attributes | key/value pairs stored in the instance state |
| class (static) method | exported module function |

Annex 7: Overall modelling and simulation approach (example from the CLEVER project)

Massive roll out of smart meters in Great-Britain

A challenge raising questions

Volume of data transfers | Associated costs | Scalability of the system

What is the most appropriate smart metering architecture?

DEFINE a Set of Experiments

Covering all questions raised by a smart metering systems roll-out

DESCRIBE

Business processes
for Smart Metering from
meter reading to advanced
DSM including future needs

CHARACTERISE

Architecture options
System structure,
telecom solutions

DEVELOP and VALIDATE

Simulation engine & tools
Simulator, models,
computing resources,
formal descriptions

RUN the Set of Experiments

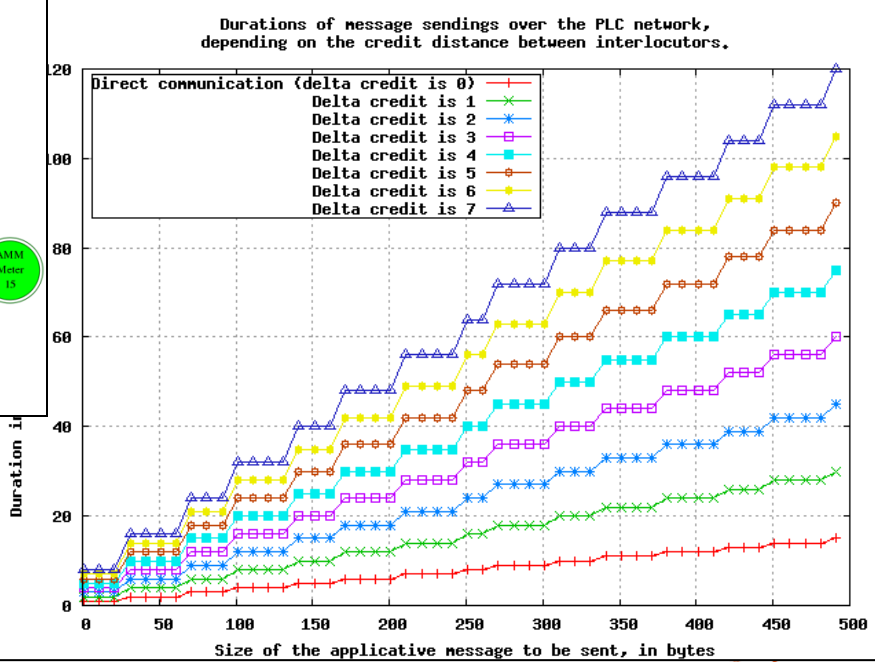
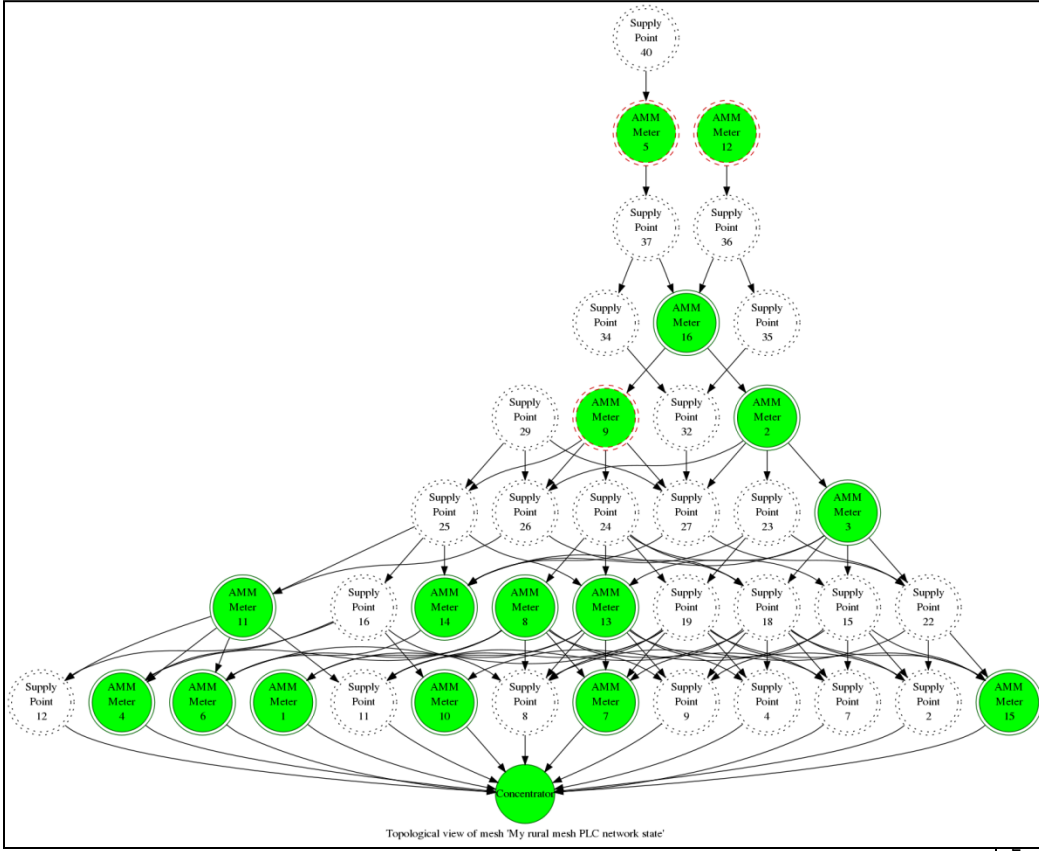
To get raw results: curves, probe indicators

ANALYSE Results

To interpret data and formulate clearly understandable conclusions

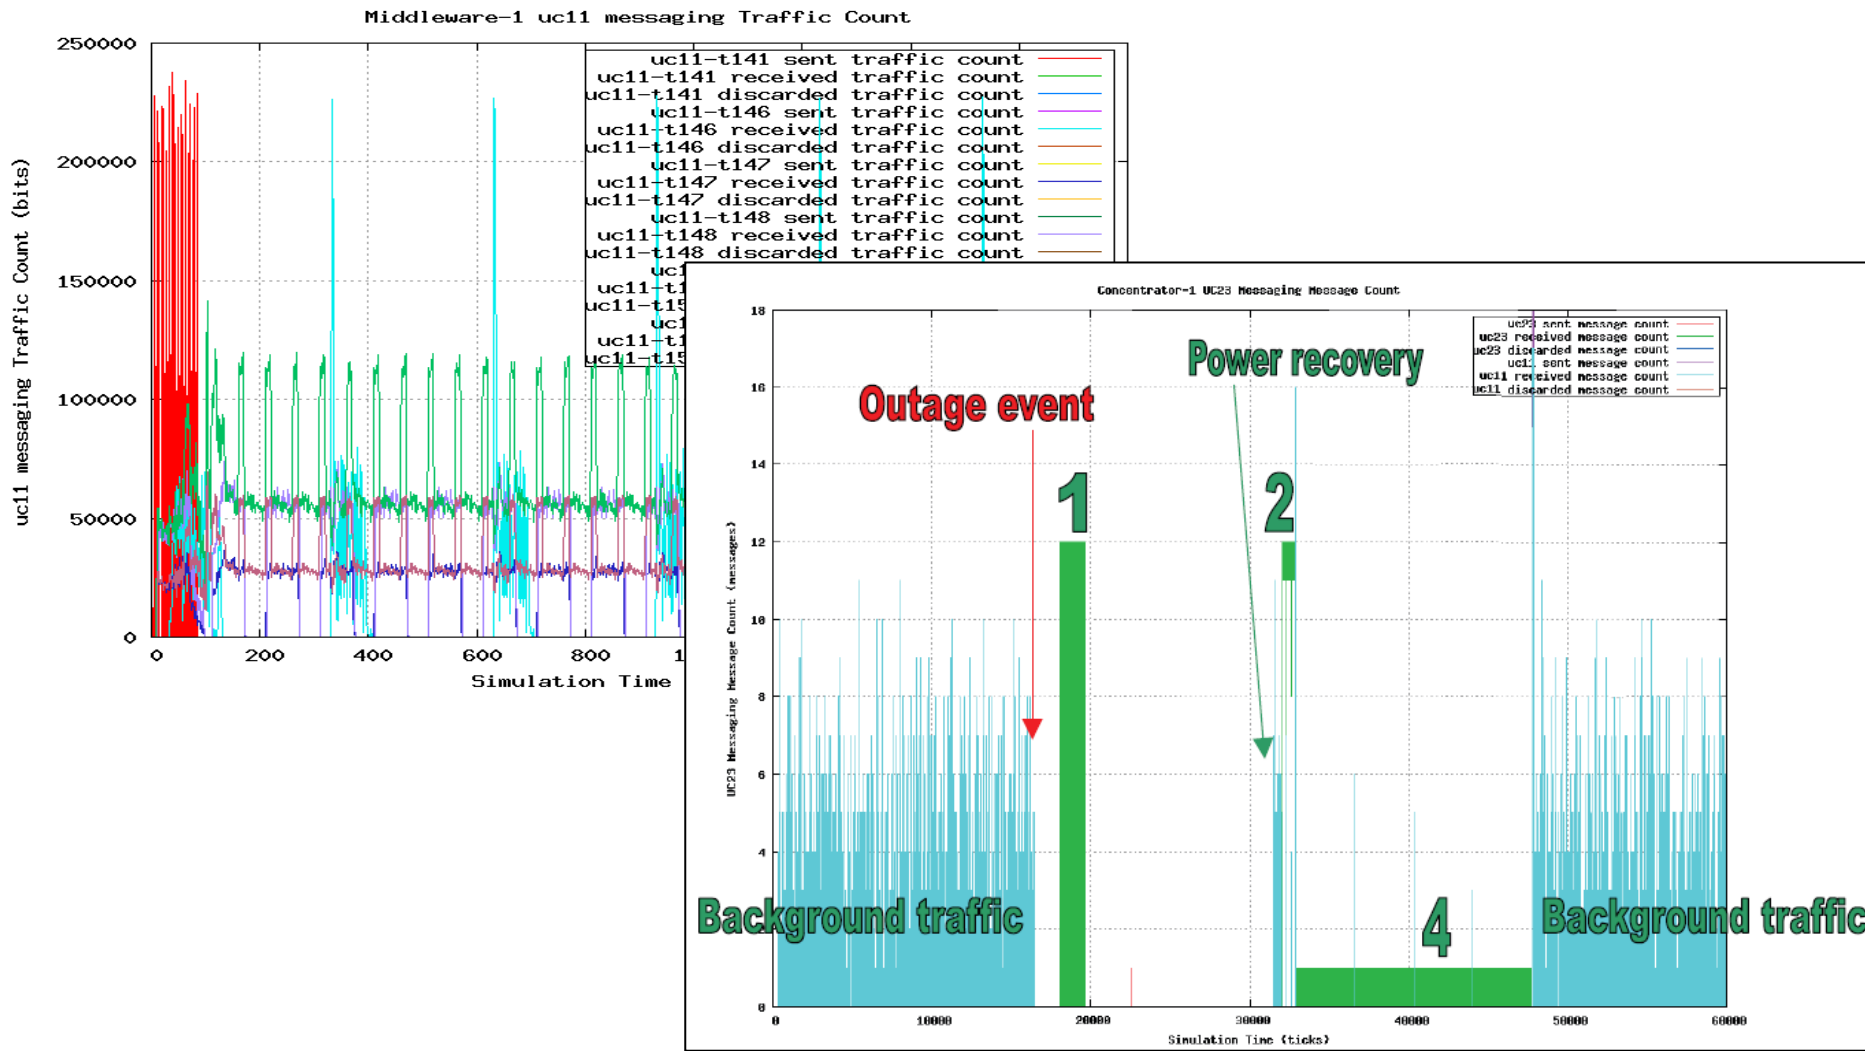
Annex 8: Examples of Sim-Diasca results & outputs (1/3)

Simulation of meters communicating with a concentrator through PLC G1
(French case, 2008)



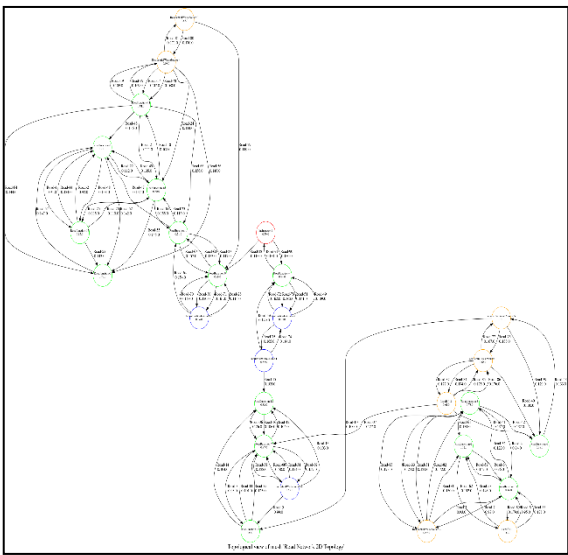
Annex 8: Examples of Sim-Diasca results & outputs (2/3)

Simulation results obtained through the CLEVER simulator: evaluating outcomes and performances depending on the functional and technical architecture of various smart metering systems (British case, 2009-2011)



Annex 8: Examples of Sim-Diasca results & outputs (3/3)

Simulation results obtained, on the left, through the RELEASE European project (large scale City-example case, here the road network of the smallest scale); on the right, internal mock-up Sustainable Cities case, devised for platform integration.



Road network (for scale 'tiny')

```

Running a Sustainable-Cities simulation from an initial state read from file 'singapore-city-structure.Init' with a timestep of 1 hour and a duration of 5 years (hence corresponding to 43830 expected timesteps).

The single specified computing host is available, using corresponding node: Sim-Diasca_Sustainable-Cities_Loading_From_File_Case-boudevil@volt.der.edf.fr.

To connect to computing nodes ['Sim-Diasca_Sustainable-Cities_Loading_From_File_Case-boudevil@volt.der.edf.fr'], use cookie '2a96f03c-da7d-44f0-9bdc-bb45224d68a9'.

Loading initial instances from:
+ file 'singapore-city-structure.Init'

All 54 initial instances have been successfully loaded from the initialisation sources, in 36.11 seconds.

-----
| Simulation Time | Tick Offset | Diasca | Real Time | Actor Count | Schedulings | Process Count |
-----
S: (not started) T: 0:D: 0|R: 21/12/2015 16:06:53|A: 0|S: 0|P: 34|
S: 1/1/2020 0:00:00 T: 0:D: 0|R: 21/12/2015 16:06:53|A: 54|S: 1|P: 93|
S: 10/2/2020 20:00:00 T: 980:D: 1|R: 21/12/2015 16:06:54|A: 54|S: 16|P: 147|
S: 21/9/2020 20:00:00 T: 6356:D: 1|R: 21/12/2015 16:06:55|A: 54|S: 16|P: 147|
S: 19/7/2021 0:00:00 T: 13560:D: 1|R: 21/12/2015 16:06:56|A: 54|S: 30|P: 147|
S: 23/5/2022 12:00:00 T: 20964:D: 0|R: 21/12/2015 16:06:57|A: 54|S: 19|P: 147|
S: 28/3/2023 12:00:00 T: 28380:D: 0|R: 21/12/2015 16:06:58|A: 54|S: 19|P: 147|
S: 18/1/2024 8:00:00 T: 35480:D: 1|R: 21/12/2015 16:06:59|A: 54|S: 16|P: 147|
S: 23/11/2024 0:00:00 T: 42912:D: 0|R: 21/12/2015 16:07:00|A: 54|S: 35|P: 147|
S: 31/12/2024 20:00:00 T: 43844:D: 0|R: 21/12/2015 16:07:00|A: 54|S: 19|P: 147|
-----

Simulation terminated successfully at simulation time: 1/1/2025 0:00:00 (tick 17750808), real time: 21/12/2015 16:07:00, after a duration of 1827 days in simulation time (43848 ticks), computed during a wall-clock duration of 6 seconds and 928 milliseconds. Simulation ran faster than the clock, with an acceleration factor of x22784757.506.

In the course of this run, a total of 21927 diascas were evaluated, corresponding to a total of 443557 scheduled instances. This corresponds to a potential average concurrency of 20.2 instances evaluated per diasca (not counting any engine-level activity).

Simulation success, result reports to be processed, collected then browsed now.
Results are available now.
(displaying 54 graphical results)
The result netcdf4 file 'dataset-from-file.nc' is available.
    
```

Console Tracker

Trace Browser

Result Browser